Supplementary Material for Robust Federated Causal Estimation (RoCE) through Calibrated High-Dimensional Models

Sinian Zhang, Yuanlin Feng, and Jue Hou

Division of Biostatistics & Health Data Science
School of Public Health, University of Minnesota, Minneapolis, Minnesota, USA

Contents and Reading Guide

The proofs follow the main text's candidate-estimation and joint-aggregation arguments. The roadmap in Section S2 links each main result to its proof; Section S7 collects alternative constructions.

S1 Method and Implementation Details 2
S1.1 Source Initialization and Fold Updates 2
S1.2 Target-Only Calibration 3
S1.3 Covariance for Joint Aggregation 5
S1.4 Outer Evaluation and Batched Communication 6
S1.5 Tuning and Numerical Stabilization 7
S2 Calibrated Candidate Expansions 9
S2.1 Setup and Notation 9
S2.2 Assumptions 11
S2.3 Source Identification 13
S2.4 Population Targets and Sparsity 14
S2.5 Calibrated Population Orthogonality 16
S2.6 A Sufficient Nuisance Oracle Bound 17
S2.7 Verification of Curvature, Localization, and Rates 20
S2.8 Second-Order Bias and Score Stability 25
S2.9 Candidate Expansion and Variance Consistency 26
S2.10 Outcome Links and Continuous Responses 28
S2.11 From Arm Means to the Treatment Contrast 29
S3 Joint Aggregation and Inference 30
S3.1 Target-Only Limit Theory 30
S3.2 Covariance Decomposition 32
S3.3 Verification of Training Moments 33
S3.4 Candidate Consistency under Stable Initial Fits 35
S3.5 Proof of the Joint Aggregation Theorem 38
S3.6 Reported Standard Errors 40
S4 Simulation Models and Comparators 43
S5 All of Us Cohort Definition and Site Composition 45
S6 Right Heart Catheterization Analysis 47
S7 Alternative Validation and Armwise Aggregation 50
S7.1 Independent Inner Validation 50
S7.2 One-Round and Two-Round Communication 52
S7.3 Connecting Inner Validation to Outer Evaluation 53
S7.4 Armwise Covariance and Objective 56
S7.5 Unbiased Candidates and Separation 57
S7.6 Oracle Aggregation Result 59
S7.7 Selection and Oracle Equivalence 59
S7.8 Squared Discrepancies and Local Bias 60

S1 Method and Implementation Details

We give the fitting, covariance, and communication details for Section 3. Initial fits supply fixed calibration multipliers; calibrated fitting and aggregation-weight learning both exclude the outer evaluation fold.

S1.1 Source Initialization and Fold Updates

We implement the source calibration losses in Section 3.2 using the primary two-layer algorithm in Section 3.4. Let $K_f \geq 3$ and use its sitewise folds $\mathcal{I}_{r,k_1}$, sizes $n_{r,J}$, and averages $\widehat{\mathbb{E}}_{r,J}$. Outer fold k_1 supplies final evaluation; inner fold $k_2 \neq k_1$ supplies a calibration loss block. Initial fitting excludes both folds. The final calibrated fits pool the inner blocks, and their fitted training scores are reused for aggregation-weight learning. Independent inner validation is a separate construction in Section S7.

Write $L(Y,u) = -Yu + \Psi(u)$, where $\Psi' = g$. Following Tan (2020), RML denotes regularized initial outcome fitting, RCAL calibrated weighting, and RWL weighted outcome fitting; IRCAL denotes initial joint-weight fitting. The subscript init marks initial coefficients; calibrated fits omit it. The outcome initialization uses the target arm and is shared across sources. Each source fits its initial joint weight using local observations and a target feature mean. The fold-specific initial fits are

$$\widehat{\alpha}_{t,a,\text{init}}^{(-k_1,-k_2)} \in \arg\min_{\alpha} \left\{ \widehat{\mathbb{E}}_{t,-k_1,-k_2}[I(A=a)L(Y,\phi^T\alpha)] + \lambda_{\alpha,\text{init}}\|\alpha\|_{1,-0} \right\},$$

$$\widehat{\gamma}_{s_j,a,\text{init}}^{(-k_1,-k_2)} \in \arg\min_{\gamma} \left\{ \widehat{\mathbb{E}}_{t,-k_1,-k_2}(\phi)^T\gamma + \widehat{\mathbb{E}}_{s_j,-k_1,-k_2}[I(A=a)e^{-\phi^T\gamma}] + \lambda_{\gamma,\text{init}}\|\gamma\|_{1,-0} \right\}.$$

Within each inner fold k_2, temporarily suppress $(-k_1,-k_2)$ on the initial coefficients and define $d_{k_2}^a(\boldsymbol{X}) = g'\{\mathcal{T}_M(\phi^T\widehat{\alpha}_{t,a,\text{init}})\}$ and $\omega_{j,k_2}^a(\boldsymbol{X}) = e^{-\mathcal{T}_M(\phi^T\widehat{\gamma}_{s_j,a,\text{init}})}$. The source's calibrated

coefficients minimize the sum of losses evaluated on the inner folds, with each site's fold contributions normalized by its own outer-training sample size:

$$\widehat{\gamma}_{s_j,a}^{(-k_1)} \in \arg\min_{\gamma} \left\{ \sum_{k_2 \neq k_1} \frac{n_{t,k_2}}{n_{t,-k_1}} \widehat{\mathbb{E}}_{t,k_2}(d_{k_2}^a \phi)^T \gamma \right.$$
$$\left. + \sum_{k_2 \neq k_1} \frac{n_{s_j,k_2}}{n_{s_j,-k_1}} \widehat{\mathbb{E}}_{s_j,k_2}\{I(A=a) e^{-\phi^T \gamma} d_{k_2}^a\} + \lambda_\gamma \|\gamma\|_{1,-0} \right\},$$
$$\widehat{\alpha}_{t,s_j}^{a,(-k_1)} \in \arg\min_{\alpha} \left\{ \sum_{k_2 \neq k_1} \frac{n_{s_j,k_2}}{n_{s_j,-k_1}} \widehat{\mathbb{E}}_{s_j,k_2}\{I(A=a)\omega_{j,k_2}^a L(Y, \phi^T \alpha)\} + \lambda_\alpha \|\alpha\|_{1,-0} \right\}.$$

The target-only fits use the same inner folds: each initial propensity fit excludes both k_1 and k_2, and each calibrated target update minimizes the fold-size-weighted loss over all $k_2 \neq k_1$; see Section S1.2. Each source initial fit also excludes the fold on which its plug-in is used. Each calibrated update incurs its slope penalty once.

Section S1.3 specifies the covariance summaries from these fits; Section S1.4 describes weight learning and outer evaluation.

S1.2 Target-Only Calibration

We specify the calibrated target-only reference in Eq. (7). Fix an arm $a \in \{0, 1\}$. Retain the main text's outcome function $m(\boldsymbol{x}; \alpha) = g\{\phi(\boldsymbol{x})^T \alpha\}$ and use the logistic propensity model $e_t^a(\boldsymbol{x}; \beta) = \text{expit}\{\phi(\boldsymbol{x})^T \beta\}$. The target-only AIPW contribution is

$$h_0(D; \alpha, \beta) = m(\boldsymbol{X}; \alpha) + \frac{I(A=a)}{e_t^a(\boldsymbol{X}; \beta)}\{Y - m(\boldsymbol{X}; \alpha)\}.$$

Its population mean is $\mu_0(\alpha, \beta) = \mathbb{E}_t\{h_0(D; \alpha, \beta)\}$. The subscript 0 identifies the target-only reference. This smooth score is used in the proof; the implemented final score uses σ_M for the propensity, which is asymptotically equivalent under the established margin and ℓ_1 rate.

Let $(\bar{\alpha}_0, \bar{\beta}_0)$ be the calibrated population coefficients, with $\bar{m}_t^a(\boldsymbol{x}) = m(\boldsymbol{x}; \bar{\alpha}_0)$ and $\bar{e}_t^a(\boldsymbol{x}) = e_t^a(\boldsymbol{x}; \bar{\beta}_0)$. If the outcome model is correct, $\bar{\alpha}_0 = \alpha^\star$ and the target propensity loss with initial $\alpha^\star$ defines $\bar{\beta}_0$. If the propensity model is correct, $\bar{\beta}_0 = \beta^\star$ and the target outcome loss with initial $\beta^\star$ defines $\bar{\alpha}_0$.

Following regularized calibrated target estimation (Tan, 2020; Smucler et al., 2019; Hou et al., 2025), write $u_\alpha = \phi^T\alpha$ and $u_\beta = \phi^T\beta$. Given initial target fits $\widehat{\alpha}_{t,a,\text{init}}$ and $\widehat{\beta}_{t,a,\text{init}}$, the calibrated losses are

$$\begin{aligned}\ell_{\text{RCAL},t}(\beta;\widehat{\alpha}_{t,a,\text{init}}) &= \widehat{\mathbb{E}}_t\left[g'\{\mathcal{T}_M(u_{\widehat{\alpha}_{t,a,\text{init}}})\}\{I(A\neq a)u_\beta + I(A=a)e^{-u_\beta}\}\right] + \lambda_\beta\|\beta\|_{1,-0},\\ \ell_{\text{RWL},t}(\alpha;\widehat{\beta}_{t,a,\text{init}}) &= \widehat{\mathbb{E}}_t\left[I(A=a)e^{-\mathcal{T}_M(u_{\widehat{\beta}_{t,a,\text{init}}})}\{-Yu_\alpha + \Psi(u_\alpha)\}\right] + \lambda_\alpha\|\alpha\|_{1,-0}.\end{aligned} \tag{S1}$$

The exponential multiplier in the outcome loss is propensity odds; the final target AIPW score uses inverse propensity. Target calibration is local and requires no intersite messages.

For the primary two-layer construction, let k_1 be the outer evaluation fold and $k_2 \neq k_1$ an inner calibration fold, with target fold notation as in the main text. The initial target outcome fit uses target arm a outside both folds. The initial propensity fit uses all target observations outside both folds with response $I(A=a)$ and logistic loss $L_{\log}(B,u) = -Bu + \log(1+e^u)$:

$$\widehat{\beta}^{(-k_1,-k_2)}_{t,a,\text{init}} \in \arg\min_\beta\left\{\widehat{\mathbb{E}}_{t,-k_1,-k_2}[L_{\log}\{I(A=a),u_\beta\}] + \lambda_{\beta,\text{init}}\|\beta\|_{1,-0}\right\}.$$

For each k_2, suppress $(-k_1,-k_2)$ on the initial fits and put $d^a_{k_2} = g'\{\mathcal{T}_M(u_{\widehat{\alpha}_{t,a,\text{init}}})\}$ and $\omega^a_{t,k_2} = e^{-\mathcal{T}_M(u_{\widehat{\beta}_{t,a,\text{init}}})}$. The calibrated target fits use only target inner folds and weight each loss block by its share of the outer-training sample:

$$\begin{aligned}\widehat{\beta}^{a,(-k_1)}_t &\in \arg\min_\beta\left\{\sum_{k_2\neq k_1}\frac{n_{t,k_2}}{n_{t,-k_1}}\widehat{\mathbb{E}}_{t,k_2}\left[d^a_{k_2}\{I(A\neq a)u_\beta + I(A=a)e^{-u_\beta}\}\right] + \lambda_\beta\|\beta\|_{1,-0}\right\},\\ \widehat{\alpha}^{a,(-k_1)}_t &\in \arg\min_\alpha\left\{\sum_{k_2\neq k_1}\frac{n_{t,k_2}}{n_{t,-k_1}}\widehat{\mathbb{E}}_{t,k_2}\{I(A=a)\omega^a_{t,k_2}L(Y,u_\alpha)\} + \lambda_\alpha\|\alpha\|_{1,-0}\right\}.\end{aligned}$$

Each plug-in is fitted outside its own calibration fold, and each penalty enters the final update once.

For candidate target functions m, e, consistency and exchangeability give

$$\mathbb{E}_t\{m + I(A=a)(Y-m)/e\} - \mu^a_t = \mathbb{E}_t\left[\frac{e-e^{a,\star}_t}{e}(m-m^{a,\star}_t)\right].$$

Thus identification holds when $m = m_t^{a,\star}$ or $e = e_t^{a,\star}$. The target expansion and its proof are given in Proposition S2. The same target score also supplies the covariance reference for joint aggregation, as detailed next.

S1.3 Covariance for Joint Aggregation

For the primary joint procedure, fix an outer evaluation fold k_1 and use the training counts $n_{r,-k_1}$ from the main text. Suppress $(-k_1)$ on the fitted functions. The target-prediction contribution is $\widehat{h}^a_{t,j,i} = \widehat{m}^a_{s_j}(\boldsymbol{X}_i)$; the target-only and source-residual contributions are

$$\begin{aligned} \widehat{h}^a_{0,i} &= \widehat{m}^a_t(\boldsymbol{X}_i) + \frac{I(A_i = a)}{\widehat{e}^a_t(\boldsymbol{X}_i)}\{Y_i - \widehat{m}^a_t(\boldsymbol{X}_i)\}, \\ \widehat{h}^a_{s_j,j,i} &= I(A_i = a)\widehat{q}^a_{s_j}(\boldsymbol{X}_i)\{Y_i - \widehat{m}^a_{s_j}(\boldsymbol{X}_i)\}. \end{aligned} \tag{S2}$$

These are the contributions underlying the candidate means in Eq. (7) and Eq. (6). Unlike the independent-validation alternative in Section S7.1, the primary procedure evaluates them on the outer training observations used for calibration. Its training-moment requirements are stated in Assumption 4.

Center each contribution within its inner calibration fold. Let $\widehat{B}^{ab}$ be the pooled cross-covariance of the target vectors $(\widehat{h}^a_{0,i}, \widehat{h}^a_{t,1,i}, \ldots, \widehat{h}^a_{t,K,i})^T$ and their arm-b counterparts, with divisor $n_{t,-k_1}$. Let $\widehat{V}^{ab}_{s_j}$ be the corresponding pooled covariance between $\widehat{h}^a_{s_j,j,i}$ and $\widehat{h}^b_{s_j,j,i}$, with divisor $n_{s_j,-k_1}$. For candidate order $(t, s_1, \ldots, s_K)$, the covariance block is

$$\widehat{\Sigma}^{ab} = \frac{\widehat{B}^{ab}}{n_{t,-k_1}} + \operatorname{diag}\left(0, \frac{\widehat{V}^{ab}_{s_1}}{n_{s_1,-k_1}}, \ldots, \frac{\widehat{V}^{ab}_{s_K}}{n_{s_K,-k_1}}\right).$$

The diagonal source term follows from independence across sources; the target block retains covariance between all candidate predictions. Writing $b(\eta^a) = (1 - \sum_j \eta^a_{s_j}, (\eta^a)^T)^T$, the variance in Eq. (10) is equivalently

$$\widehat{V}^{(-k_1)}_\tau(\eta) = b(\eta^1)^T\widehat{\Sigma}^{11}b(\eta^1) + b(\eta^0)^T\widehat{\Sigma}^{00}b(\eta^0) - 2b(\eta^1)^T\widehat{\Sigma}^{10}b(\eta^0).$$

In particular, the discrepancy standard error satisfies $(\widehat{s}^a_{s_j})^2 = \widehat{\Sigma}^{aa}_{jj} + \widehat{\Sigma}^{aa}_{00} - 2\widehat{\Sigma}^{aa}_{0j}$. Candidate means use the pooled site averages before discrepancies and penalties are computed. The

full stacked two-arm covariance matrix is positive semidefinite when assembled from these common sitewise score matrices.

Disjoint arm indicators do not imply zero covariance after centering. At a source, the two uncentered residual contributions have product zero, so their covariance is minus the product of their means; it vanishes only when at least one residual mean is zero. At the target, even the true AIPW scores have cross-arm covariance $\mathrm{Cov}_t\{m_t^{1,\star}(\boldsymbol{X}), m_t^{0,\star}(\boldsymbol{X})\}$. Both facts explain why the joint objective retains cross-arm moments.

An equivalent quadratic implementation stacks the treated-arm discrepancies and the negatives of the control-arm discrepancies:

$$(\widehat{d}^1_{s_1}, \ldots, \widehat{d}^1_{s_K}, -\widehat{d}^0_{s_1}, \ldots, -\widehat{d}^0_{s_K})^T.$$

Its joint variance is the target-only contrast variance, plus twice the covariance of that contrast with the weighted increments, plus the variance of the weighted increments. Multiplication by $n = n_{-k_1}$ gives the quadratic coefficients in Assumption 5; the minus signs on control-arm increments are essential.

S1.4 Outer Evaluation and Batched Communication

The source and target fits above supply the candidate means in Eq. (6) and Eq. (7). For each outer fold k_1, use only its complement to compute those means and the covariance blocks in Section S1.3. Pool the training moments within each site before forming discrepancies and solving the joint weight objective in Eq. (11). The calibrated models and learned weights then determine the held-out site contributions in Eq. (9). Pooling those contributions with each site's full sample size gives Eq. (12). This order separates final evaluation from fitting and weight learning.

The inner folds separate each initial fit from the observations contributing its calibration loss. The final calibrated fits nevertheless use the outer training sample on which aggregation

weights are learned. The training-score control stated in Assumption 4 is verified by Proposition S3. The independent-inner-validation alternative is developed separately in Section S7.1.

With target outcome initialization, the fitting tasks for all folds can be declared in advance and batched into one target broadcast and one source return. The target sends initial outcome coefficients, initial-training feature means, and calibration-fold derivative moments. Each source returns calibrated outcome coefficients, residual means, second moments, cross-arm products, and counts for the specified training and evaluation tasks. The target computes prediction covariances and joint aggregation weights locally; evaluation summaries do not enter weight learning. This is the protocol in figure 1. Its first- and second-moment messages support point estimation and the leading variance. Additional products for the optional local weight-sensitivity diagnostic, which is not included in the primary standard errors, are described in Section S3.6.

S1.5 Tuning and Numerical Stabilization

For the sufficient theoretical regime below, use a common feature basis, intercept convention, and clipping radius across folds and sites. The intercept is unpenalized; all slope coefficients receive an ℓ_1 penalty. Validation, if used, is confined to the corresponding training sample and matches the fitted loss. The inference theorem additionally requires the high-probability gradient event and restricted curvature below; it does not assert that a validation-selected penalty automatically satisfies them.

We use the clipping convention from Section 3.2: $\mathcal{T}_M(u) = \operatorname{sign}(u)\min(|u|, 2M)$ and $\sigma_M(u) = \operatorname{expit}\{\mathcal{T}_M(u)\}$. The population bound is M and the numerical predictor radius is $2M$, with $M > 0$ fixed. The following prescriptions specify where clipping enters the fitting sequence.

(a) In the source joint-weight loss, use the identical factor $g'\{\mathcal{T}_M(\phi^T\widehat{\alpha}_{t,a,\text{init}})\}$ for the target linear moment and the source exponential term.

(b) In the source outcome loss, use $I(A=a)\exp\{-\mathcal{T}_M(\phi^T\widehat{\gamma}_{s_j,a,\text{init}})\}$. In the target outcome loss, use $I(A=a)\exp\{-\mathcal{T}_M(\phi^T\widehat{\beta}_{t,a,\text{init}})\}$.

(c) In the target propensity loss, the single factor $g'\{\mathcal{T}_M(\phi^T\widehat{\alpha}_{t,a,\text{init}})\}$ multiplies the entire bracket $(1-I(A=a))\phi^T\beta + I(A=a)\exp(-\phi^T\beta)$.

(d) Evaluate final outcome means as $g(\phi^T\widehat{\alpha})$, final source weights as $\exp(-\phi^T\widehat{\gamma})$, and calibrated target propensities as $\sigma_M(\phi^T\widehat{\beta}_t^a)$. Thus target probabilities lie in $[\text{expit}(-2M), \text{expit}(2M)]$. No additional probability clipping is imposed.

The final target clipping is asymptotically inactive under the margin and ℓ_1 consistency conditions below, including when the final propensity model is misspecified. Initial clipping need not be inactive for a misspecified model. This is the role of fixed predictor clipping in the casewise proof; an arbitrarily selected radius alone does not establish overlap or inference. The construction uses the clipping principle of Hou et al. (2025, eqs. (11)–(13), Assumption 4); the transported source score requires its own two-sample balance argument.

For a sample of total size n used to learn aggregation weights, $c_n = \sqrt{\log(n \vee 3)}$ is one choice satisfying the theorem's cutoff conditions and gives penalty $(\widehat{t}_j/c_n - 1)_+$. Here n excludes the final evaluation fold. The joint covariance is assembled from target and source empirical covariance blocks; a ridge, if used, must vanish asymptotically. If an unscaled discrepancy variance is floored, $n\epsilon_{V,n} \to 0$ is sufficient to preserve its stated limiting scale. The numerical analyses use a fixed Wald cutoff of 2 and the application-specific predictor radii reported in the main text, which can differ across sites. These settings, including the fixed variance floor in Section S3.6, are distinct from the common-radius, growing-cutoff sufficient regime of the theory.

S2 Calibrated Candidate Expansions

This section proves source identification and calibration from Section 3.1 and Section 3.2, followed by the source expansion in Section 4.1. The argument proceeds from population targets and orthogonality to nuisance rates, second-order remainders, and cross-fitted inference.

Main-text results are restated under their original numbers before their proofs. Only additional supplementary results receive S-numbers. The table gives the proof locations and ingredients.

Main result	Proof ingredients	Proof page
Lemma 1	Direct identification argument in Section S2.3	13
Lemma 2	Population stationarity in Section S2.5	16
Lemma 3	Sitewise covariance decomposition in Section S1.3	32
Lemma 4	Lemmas S3 and S4: nuisance rates and second-order bounds	26
Theorem 1	Source and target expansions, stable biased candidates, and Proposition S3	38

S2.1 Setup and Notation

Fix a source s, target t, and arm $a \in \{0, 1\}$. As in the main text, n_r is the site sample size, $N = \sum_r n_r$ is the full network size, and $n_r/N \to \rho_r^{\text{all}} > 0$. The pair size $n_{t,s} = n_t + n_s$ and all fixed-fold fitting sample sizes are therefore proportional to N.

For coefficient vectors $\alpha, \gamma \in \mathbb{R}^d$, the two working functions are

$$m(\boldsymbol{x}; \alpha) = g\{\phi(\boldsymbol{x})^T \alpha\}, \qquad q_s^a(\boldsymbol{x}; \gamma) = \exp\{-\phi(\boldsymbol{x})^T \gamma\}.$$

The first is an outcome mean; the second is a joint weight for source observations in arm a. The common feature vector $\phi(\boldsymbol{x})$ includes an intercept, while p denotes the raw covariate dimension. The ideal joint weight is

$$q_s^{a,\star}(\boldsymbol{x}) = \frac{f_t^\star(\boldsymbol{x})}{f_s^\star(\boldsymbol{x}) e_s^{a,\star}(\boldsymbol{x})}.$$

Stack the coefficients as $\vartheta = (\alpha^T, \gamma^T)^T \in \mathbb{R}^{2d}$. For a single record $D = (R, \boldsymbol{X}, A, Y)$, the target prediction is $h_t(D; \vartheta) = m(\boldsymbol{X}; \alpha)$ and the source residual contribution is

$$h_s(D; \vartheta) = I(A = a) q_s^a(\boldsymbol{X}; \gamma)\{Y - m(\boldsymbol{X}; \alpha)\}.$$

The subscript indicates the site at which each contribution is averaged. Thus h_t contains only the outcome prediction; the target-only AIPW contribution is h_0 from Section S1.2.

Their population mean is the same source-assisted mean defined in Eq. (2), with the source and arm fixed here:

$$\mu_s^a(\vartheta) = \mu_s^a(\alpha, \gamma) = \mathbb{E}_t\{h_t(D; \vartheta)\} + \mathbb{E}_s\{h_s(D; \vartheta)\}.$$

The stacked argument is used for joint Taylor expansions; the population mean retains the source and arm indices of the main text.

Let $\bar{\vartheta} = (\bar{\alpha}^T, \bar{\gamma}^T)^T$ collect the calibrated population coefficients. The corresponding functions are $\bar{m}_s^a(\boldsymbol{x}) = m(\boldsymbol{x}; \bar{\alpha})$ and $\bar{q}_s^a(\boldsymbol{x}) = q_s^a(\boldsymbol{x}; \bar{\gamma})$, as in Section 4.1. Their source-assisted population mean is $\mu_s^a(\bar{\vartheta})$. The subscript init also identifies initial population minimizers and support sizes. The two correct-model cases defining the calibrated limits are specified in Assumption S3.

For the asymptotic expansion, center each contribution under its own site's distribution:

$$\zeta_{t,s}^a = h_t(D; \bar{\vartheta}) - \mathbb{E}_t\{h_t(D; \bar{\vartheta})\}, \qquad \xi_{t,s}^a = h_s(D; \bar{\vartheta}) - \mathbb{E}_s\{h_s(D; \bar{\vartheta})\}.$$

These are the scores in Section 4.1. We write $\widehat{\mu}_s^a$ for the cross-fitted source-assisted mean. Its variance on the main text's $\sqrt{N}$ scale is $V_{t,s}^{a,*} = (\rho_t^{\text{all}})^{-1} \operatorname{Var}_t(\zeta_{t,s}^a) + (\rho_s^{\text{all}})^{-1} \operatorname{Var}_s(\xi_{t,s}^a)$. Finally, write s_c, s_w for the support sizes of the correct and other calibrated coefficients, including the intercept even when it is zero, and $s_{c,\text{init}}$ for the correct initial support size. Put $s_{\text{eff}} = s_c + s_w + s_{c,\text{init}}$. When both models are correct, include both initial support sizes; this only changes a fixed factor when the initial and calibrated models share the same basis. Coefficient norms are denoted by $\|\cdot\|_q$ and function norms by $\|\cdot\|_{L_2(P_r)}$.

S2.2 Assumptions

We retain the causal identification conditions in Assumption 1. The first two assumptions below give the population-design, tail, overlap, and predictor-margin conditions used in Assumption 2. The third specifies the correct-model branches and population supports underlying Assumption 3; its sparsity rate is imposed when the inference result is applied.

The later oracle-bound event in Assumption S4 is verified from these primitive conditions by Lemma S3. Training moments and source separation are treated in Section S3.

Assumption S1 (Sampling and population design)**.** Observations are i.i.d. within sites and independent across sites. Site fractions have positive limits, and the numbers of sites and all cross-fitting folds are fixed. Fold memberships are balanced and chosen independently of the observed records, as in Section 3.4. At each site used for fitting, $\|\phi\|_\infty \le C_\phi$. For each fitted arm, let $G_{r,a} = \mathbb{E}_r\{I(A=a)\phi\phi^T\}$ and assume

$$\lambda_{\min}(G_{r,a}) \ge \kappa_G > 0, \qquad \lambda_{\max}\{\mathbb{E}_r(\phi\phi^T)\} \le C_\Sigma. \tag{S3}$$

Constants do not depend on d, N. The common, prespecified feature basis includes an intercept and has no exact redundant columns. The lower bound is a condition on population matrices; empirical matrices may be singular when d exceeds the sample size.

Assumption S2 (Link, tails, overlap, and predictor margins)**.** Use the observed conditional means $m_r^{a,\star}(\boldsymbol{x}) = \mathbb{E}_r(Y \mid \boldsymbol{X} = \boldsymbol{x}, A = a)$ from Section 3.1.

(a) The outcome link $g = \Psi'$ is twice continuously differentiable on $\mathbb{R}$, with $g' > 0$. For a fixed $M > 0$, on $[-2M, 2M]$,

$$0 < m_g \le g'(u) \le C_g, \qquad |g(u)| + |g''(u)| \le C_g.$$

Treatment propensities use $\sigma(u) = \mathrm{expit}(u)$ regardless of g. Use $\mathcal{T}_M$ and σ_M as specified in Section 3.2.

(b) For some fixed $B < \infty$, $g(-B) \le m_r^{a,\star}(\boldsymbol{X}) \le g(B)$ almost surely. For $\epsilon_{r,a} = Y - m_r^{a,\star}(\boldsymbol{X})$, there are fixed $v_Y, b_Y > 0$ such that, for $|z| < 1/b_Y$,

$$\mathbb{E}_r\{e^{z\epsilon_{r,a}} \mid \boldsymbol{X}, A = a\} \le \exp\left\{\frac{v_Y z^2}{2(1 - b_Y|z|)}\right\}. \tag{S4}$$

This is a conditional tail condition, not a requirement that Y be bounded or Gaussian.

(c) Target propensities satisfy $\varepsilon \le e_t^{a,\star}(\boldsymbol{X}) \le 1 - \varepsilon$. On the common target–source support the ideal joint weights satisfy $0 < q_- \le q_s^{a,\star}(\boldsymbol{X}) \le q_+ < \infty$. This quantitative overlap condition does not require $q_s^{a,\star}$ to belong to the exponential working family.

(d) Calibrated population predictors and the true predictors of the correctly specified initial models are bounded in absolute value by M on the relevant supports. No bound or convergence requirement is imposed on a misspecified initial fit for the correct-branch inference result: it enters the updates through bounded clipped functions. Sitewise limiting score covariances converge, with positive variances for the contrasts to be studentized.

Section S2.10 discusses the covered outcome links and response distributions.

The conditions imply dimension-free local analytic bounds without bounding realized outcomes. In a convex neighborhood $\mathcal{N}$ in which fitted linear predictors have magnitude at most $2M$, $q_s^a(\boldsymbol{X}; \gamma) \le e^{2M}$, the conditional means and residual second moments are bounded, and

$$\begin{aligned} \sup_{\vartheta \in \mathcal{N}} \|\nabla^2 \mu_s^a(\vartheta)\|_{\mathrm{op}} &\le C, \\ \sum_{r \in \{t,s\}} \mathbb{E}_r\{h_r(D;\vartheta) - h_r(D;\bar{\vartheta})\}^2 &\le C\|\vartheta - \bar{\vartheta}\|_2^2. \end{aligned} \tag{S5}$$

For the Hessian, condition on $(\boldsymbol{X}, A)$ to replace Y by $m_r^{a,\star}(\boldsymbol{X})$; every block is then an expectation of $\phi\phi^T$ times a bounded scalar. For score stability, squared residual factors are controlled by their conditional second moments, and $\mathbb{E}_r(\phi^T\Delta)^2 \le C_\Sigma\|\Delta\|_2^2$. The same argument applies to the target score using $1/e_t^a(\boldsymbol{X}; \beta) = 1 + e^{-\phi^T\beta}$. The conditional

tail bound supplies uniformly bounded score moments of order $2+\delta$. The empirical fits' membership in $\mathcal{N}$ is proved in Lemma S3, rather than assumed here.

Assumption S3 (Correct-model branches and population supports)**.** The source is conditionally mean-transportable and either $m(\boldsymbol{X};\alpha^\star) = m_t^{a,\star}(\boldsymbol{X})$ or $q_s^a(\boldsymbol{X};\gamma^\star) = q_s^{a,\star}(\boldsymbol{X}) = f_t^\star/(f_s^\star e_s^{a,\star})$. In the outcome-correct branch, set $\bar{\alpha} = \alpha^\star$ and define $\bar{\gamma}$ by the population source joint-weight loss with initial outcome $\alpha^\star$. In the weight-correct branch, set $\bar{\gamma} = \gamma^\star$ and define $\bar{\alpha}$ by the population source outcome loss with initial joint-weight fit $\gamma^\star$. The supports of these two calibrated coefficients and the correct initial coefficient have total size s_{eff}, counting intercepts. When both models are correct, include both initial supports. All fitting is confined to the appropriate training sample, using penalties specified below. For the target, replace the source weight branch by correctness of its logistic propensity model and use the corresponding target population losses.

S2.3 Source Identification

Lemma 1 (restated)**.** *For a transportable source and coefficient vectors α, γ with finite expectations, Assumption 1 gives*

$$\mu_{s_j}^a(\alpha,\gamma) - \mu_t^a = \mathbb{E}_t\left[\left(1 - \frac{q_{s_j}^a(\boldsymbol{X};\gamma)}{q_{s_j}^{a,\star}(\boldsymbol{X})}\right)\left(m(\boldsymbol{X};\alpha) - m_t^{a,\star}(\boldsymbol{X})\right)\right].$$

Thus $\mu_{s_j}^a(\alpha,\gamma) = \mu_t^a$ if $m(\boldsymbol{X};\alpha) = m_t^{a,\star}(\boldsymbol{X})$ or $q_{s_j}^a(\boldsymbol{X};\gamma) = q_{s_j}^{a,\star}(\boldsymbol{X})$.

Proof of Lemma 1. Conditional mean transportability gives $\mathbb{E}_s(Y \mid \boldsymbol{X}, A = a) = m_t^{a,\star}(\boldsymbol{X})$. Conditioning the source residual in Eq. (2) therefore replaces Y by this mean. The identity $f_s^\star e_s^{a,\star} q_s^{a,\star} = f_t^\star$ then converts the source expectation to a target expectation with multiplier $q_s^a(\boldsymbol{X};\gamma)/q_s^{a,\star}(\boldsymbol{X})$. Combining it with the target prediction and subtracting $\mu_t^a = \mathbb{E}_t\{m_t^{a,\star}(\boldsymbol{X})\}$ gives the stated identity. Either correct working model makes the resulting bias zero. □

S2.4 Population Targets and Sparsity

To use stationarity at the calibrated limits, we first establish finite, unique population minimizers. We then give a sufficient sparse class under misspecification.

Proposition S1 (Finite population targets and a sparse sufficient class)**.** *Under Assumptions S1 and S2(a)–(c), every initial population loss and every calibrated population loss with a fixed clipped plug-in has a unique finite minimizer for each d. The same holds for untruncated population plug-ins whose predictors have a finite bound.*

For a sufficient sparsity construction, partition $\phi = (Z^T, W^T)^T$, with the intercept in $Z \in \mathbb{R}^q$. Suppose, at each site and for both arms,

$$\mathbb{E}_r(W \mid Z, A = a) = 0, \qquad m_r^{a,\star}(\boldsymbol{X}) = h_{r,a}(Z). \tag{S6}$$

Then all initial population coefficients, and the calibrated population coefficients using those initial population minimizers as plug-ins, have zero W coordinates, even when the working outcome or weighting model is misspecified. If q and the joint law of (Z, A, Y, R) are fixed as inactive coordinates are added, their nonzero coefficients remain fixed. A common finite M can then be chosen to bound the population predictors of all initial and calibrated fits. Thus this class verifies both sparse population targets and the predictor-margin requirement while allowing $d \gg N$.

Proof of Proposition S1. For the outcome loss, conditioning on $(\boldsymbol{X}, A = a)$ gives $\Psi(u) - m_r^{a,\star}(\boldsymbol{X})u$. The interior mean bound in Assumption S2(b) and strict monotonicity of g imply

$$\Psi(u) - \nu u \geq c|u| - C \quad \text{uniformly for } \nu \in [g(-B), g(B)].$$

For example, for $u \geq B+1$ its derivative is at least $g(B+1) - g(B) > 0$, and the negative tail is analogous. Positive bounded calibration weights preserve this lower bound. For a source joint-weight with multiplier $c(\boldsymbol{X}) > 0$, change measure to obtain

$$L_\gamma(\gamma) = \mathbb{E}_s[I(A=a)c(\boldsymbol{X})\{q_s^{a,\star}(\boldsymbol{X})\phi^T\gamma + e^{-\phi^T\gamma}\}].$$

Uniformly for $w \in [q_-, q_+]$, $wu + e^{-u} \geq c|u| - C$. Target propensity calibration instead has conditional loss $c(\boldsymbol{X})\{(1-e_t^{a,\star})u + e_t^{a,\star}e^{-u}\}$, with the same bound by overlap; the initial logistic likelihood has the corresponding interior-mean bound. All clipped multipliers have fixed positive lower and finite upper bounds. Moreover,

$$\mathbb{E}_r\{I(A=a)|\phi^T v|\} \geq \frac{\mathbb{E}_r\{I(A=a)(\phi^T v)^2\}}{C_\phi \|v\|_1} \geq \frac{\kappa_G}{C_\phi\sqrt{d}}\|v\|_2$$

for $v \neq 0$. Hence each loss is coercive for each finite d and attains its minimum. Strictly positive scalar Hessian multipliers and positive population Gram matrices make each loss strictly convex, giving uniqueness and population stationarity. These existence arguments do not claim a dimension-free coefficient bound on their own.

For sparsity, first minimize each initial loss over coefficients of Z only and set the W coefficients to zero. The Z gradient vanishes by restricted minimization. Outcome-loss W gradients vanish because the conditional mean is $h_{r,a}(Z)$ and Eq. (S6) holds. Source joint-weight gradients in W vanish in both target and source terms separately. For the target propensity likelihood, $\mathbb{E}_t(W \mid Z) = 0$ and $\mathbb{E}_t(I(A=a)W \mid Z) = 0$ give the same conclusion. Uniqueness identifies these as the full population minima.

The resulting initial predictors are functions of Z. Therefore clipped outcome derivatives and clipped exponential initial weights are also functions of Z, and the identical gradient argument applies to both calibrated losses at the source and target. This proves sparsity for the entire population fitting sequence without equating initial and calibrated coefficients.

If the active law and q are fixed, the reduced initial coefficients are finite constants. Construct the reduced calibrated population fits using these untruncated initial predictors;

the predictors are bounded because Z is bounded, so the preceding existence proof still applies. There are finitely many resulting coefficients. Choose M larger than all their predictor bounds. Initial population clipping is then inactive and these are also the targets of the clipped procedure with that M. This supplies a nonempty model class with a common margin; it is not a data-driven estimator of the unknown bound. □

For example, bounded inactive features independent of (Z, A, Y, R) and centered at zero satisfy Eq. (S6). The outcome mean may be a nonlinear function such as $h_{r,a}(Z) = g(b_0 + b_1 Z_1 + b_2 Z_1^2)$ while the working active basis contains only $(1, Z_1)$, so its projection remains sparse despite misspecification. Site and treatment mechanisms depending on Z can likewise have misspecified working joint-weight models. Correctness of one nuisance, when required for inference, is imposed separately. With growing active dimension, one must still impose the stated sparsity rate and a uniform predictor bound; this proposition does not make every misspecified projection sparse.

S2.5 Calibrated Population Orthogonality

Stationarity at the minimizers in Proposition S1 gives the calibration result.

Lemma 2 (restated)**.** *For a transportable source and one arm, let $(\bar{\alpha}, \bar{\gamma})$ be the identifiable calibrated population coefficients satisfying the predictor bound in Section 3.2. If the outcome model is correct at $\alpha^\star$, $\bar{\alpha} = \alpha^\star$; if the joint weight model is correct at $\gamma^\star$, $\bar{\gamma} = \gamma^\star$. In either case, the other coefficient minimizes its paired population calibration loss with the correct initial coefficient. Then $\mu^a_{s_j}(\bar{\alpha}, \bar{\gamma}) = \mu^a_t$, and both gradients in Eq. (3) vanish at $(\bar{\alpha}, \bar{\gamma})$.*

Proof of Lemma 2. We verify that both derivatives in Eq. (3) vanish at the calibrated population coefficients.

If the outcome model is correct, conditional mean transportability identifies $\bar{\alpha}_{\text{init}} = \alpha^\star$ for target initialization (one round) or source initialization (two rounds). Any positive clipped initial weight preserves this outcome population minimizer, so $\bar{\alpha} = \bar{\alpha}_{\text{init}} = \alpha^\star$. Conditional residual centering gives $\nabla_\gamma \mu_s^a = 0$ for every calibrated joint weight. Because $|\phi^T \alpha^\star| \le M$, the derivative multiplier in the limiting calibrated joint-weight loss is $g'(\phi^T \alpha^\star)$, so its stationarity gives $\nabla_\alpha \mu_s^a = 0$.

If instead the joint weight model is correct, the initial balance identity identifies $\bar{\gamma}_{\text{init}} = \gamma^\star$. The same identity multiplied by any common clipped outcome derivative identifies $\bar{\gamma} = \gamma^\star$ in the calibrated joint-weight loss. True weighting gives $\nabla_\alpha \mu_s^a = 0$ for every calibrated outcome coefficient. The initial true joint weight is inside the predictor margin, so the limiting outcome calibration is weighted by $q_s^a(\boldsymbol{X}; \gamma^\star)$; its stationarity gives $\nabla_\gamma \mu_s^a = 0$. The misspecified coefficients may differ across stages in either branch. In either branch, Lemma 1 also gives $\mu_s^a(\bar{\alpha}, \bar{\gamma}) = \mu_t^a$, completing the calibration result. □

S2.6 A Sufficient Nuisance Oracle Bound

Following the cross-fitted argument of Hou et al. (2025, Lemma A21), we first bound nuisance error in terms of empirical fluctuations and propagated initialization error. Lemma S3 verifies the required event and gives the final rates.

Fix an outer fold k_1. Let L_N be the unpenalized part of one pooled calibration objective in Sections S1.1 and S1.2, with one common coefficient vector and the prescribed sitewise fold weights $w_{r,k_2} = n_{r,k_2}/n_{r,-k_1}$. Each site's weights sum to one. For a population coefficient $\bar{v}$, let S contain its support and the intercept, $s_v = |S|$, and $\widehat{v} \in \arg\min_v \{L_N(v) + \lambda \|v\|_{1,-0}\}$. Write $r_{c,\text{init}}$ for the largest ℓ_2 error of the correct initial fit across the fixed inner folds.

Assumption S4 (Oracle-bound event)**.** Write $\nabla L_N(\bar{v}) = A_N + B_N$. For each inner block, condition only on its own initial-training data, center its empirical gradient, and then sum

using the original sitewise weights. This gives A_N; the weighted sum of the corresponding conditional expectations gives B_N. On the event under consideration, $\|A_N\|_\infty \le \lambda/2$ and $\|B_N\|_2 \le b_N$. On the expanded cone

$$\mathcal{C}_N = \left\{ \Delta : \|\Delta_{S^c}\|_1 \le 3\|\Delta_S\|_1 + \frac{2b_N}{\lambda}\|\Delta\|_2 \right\},$$

the loss has restricted curvature

$$L_N(\bar{v}+\Delta) - L_N(\bar{v}) - \nabla L_N(\bar{v})^T\Delta \ge \frac{\kappa}{2}\|\Delta\|_2^2$$

for a fixed $\kappa > 0$. If this inequality is local, it must hold on the relevant segments and the minimizer must be localized in that region. For a correct initial fit, the same conditions apply with $B_N = 0$.

Lemma S1 (Nuisance rate with initial-error propagation)**.** *On the preceding event,*

$$\kappa\|\widehat{v} - \bar{v}\|_2 \le 3\sqrt{s_v}\lambda + 2b_N,$$
$$\kappa\|\widehat{v} - \bar{v}\|_1 \le (4\sqrt{s_v} + 2b_N/\lambda)(3\sqrt{s_v}\lambda + 2b_N).$$

Under Assumptions S1 to S3, the pooled calibration updates admit $b_N \le C/N$ for a correctly specified model and $b_N = O_p(r_{c,\text{init}})$ for the other model.

Proof of Lemma S1. Write $\Delta = \widehat{v} - \bar{v}$. The basic inequality for an unpenalized intercept and penalized slopes gives

$$L_N(\bar{v}+\Delta) - L_N(\bar{v}) \le \lambda(\|\Delta_S\|_1 - \|\Delta_{S^c}\|_1).$$

Convexity and the gradient decomposition lower-bound the left side by $-\lambda\|\Delta\|_1/2 - b_N\|\Delta\|_2$. This gives the expanded cone. Restricted curvature then yields

$$\begin{aligned}\frac{\kappa}{2}\|\Delta\|_2^2 &\le \frac{3\lambda}{2}\|\Delta_S\|_1 - \frac{\lambda}{2}\|\Delta_{S^c}\|_1 + b_N\|\Delta\|_2 \\ &\le \left(\frac{3\lambda\sqrt{s_v}}{2} + b_N\right)\|\Delta\|_2.\end{aligned}$$

Division gives the ℓ_2 bound; the expanded cone gives the ℓ_1 bound. Including the intercept in S makes the displayed upper bounds valid without penalizing it. This is the basic-inequality argument of Negahban et al. (2012) with an explicit propagated gradient.

To control B_N, let $\mathcal{F}_{k_2}$ contain the initial-training data outside k_1, k_2. The observations in calibration fold k_2 are independent of this sigma-field. A correct outcome gives a conditionally mean-zero residual for every fixed clipped initial weight, so every source or target outcome block has zero conditional gradient. At the target, a correct propensity gives

$$\mathbb{E}_t\Big[d^a_{k_2}\{I(A \neq a) - I(A = a)e^{-\phi^T\beta^\star}\}\phi \mid \mathcal{F}_{k_2}\Big] = 0.$$

Both target terms use the same sitewise fold fraction, so their pooled conditional gradient is also zero.

For a correct source joint weight, the balance identity applies to the same derivative multiplier $d^a_{k_2}$ on both populations. Its pooled conditional gradient is

$$B_N = \sum_{k_2 \neq k_1} (w_{t,k_2} - w_{s,k_2}) m_{k_2}, \qquad m_{k_2} = \mathbb{E}_t(d^a_{k_2}\phi \mid \mathcal{F}_{k_2}). \tag{S7}$$

The bounded derivative multiplier and population Gram upper bound give $\|m_{k_2}\|_2 \leq C$ by Cauchy–Schwarz, uniformly in d. Balanced sitewise folds have fractions differing by $O(N^{-1})$. With fixed folds and positive site fractions, Eq. (S7) therefore has norm at most C/N. Exact equality of target and source fractions makes this term zero, but is not required for the rates.

For the misspecified update, compare its conditional population gradient with the stationary gradient obtained from the true correct initial coefficient. Both $g' \circ \mathcal{T}_M$ and $\exp(-\mathcal{T}_M)$ are bounded and Lipschitz. For any unit vector v and initial error Δ_{init}, each gradient difference is bounded by sums of terms of the form

$$C\mathbb{E}_r\{|v^T\phi|\,|\phi^T\Delta_{\text{init}}|\} \leq C\{\mathbb{E}_r(v^T\phi)^2\}^{1/2}\{\mathbb{E}_r(\phi^T\Delta_{\text{init}})^2\}^{1/2} \leq C\|\Delta_{\text{init}}\|_2.$$

For outcome-gradient propagation, conditioning on $(\boldsymbol{X}, A)$ replaces the residual by its bounded conditional mean; the final exponential factors are bounded by the population

margin. The stationary comparison uses one common population plug-in. Since each site's weights sum to one, its target and source terms cancel after pooling even if their fold fractions differ. Taking the supremum over v and summing the fixed number of weighted blocks gives $b_N = O_p(r_{c,\mathrm{init}})$. The correct source joint-weight remainder is absorbed by $\sqrt{s_c}\lambda$ when $\lambda \gtrsim N^{-1/2}$, since $s_c \geq 1$. These bounds supply the initialization and calibration steps in Lemma S3. □

S2.7 Verification of Curvature, Localization, and Rates

We verify the event in Lemma S1 by establishing finite initial fits, local curvature, and localization of calibrated minimizers. Let F be the fixed number of site–arm–fold empirical blocks used by the fitting procedure, including fits internal to the aggregation training samples, and let n be their smallest sample size. All such block sizes are proportional to the relevant total sample size. For $\delta \in (0,1)$, put

$$t_\delta = \log(64Fd^2/\delta), \qquad a_n = \sqrt{t_\delta/n} + t_\delta/n, \qquad \lambda = La_n. \tag{S8}$$

Different penalties may be used within a fixed constant factor of λ, with their lower multiplier sufficiently large. Define $\sigma(u) = \mathrm{expit}(u)$ and

$$h_* = e^{-2M}\min\{1, m_g, \sigma'(2M)\} > 0, \qquad \kappa_* = h_*\kappa_G/2, \qquad r_0 = M/C_\phi.$$

Lemma S2 (Finite empirical initial fits)**.** *Under Assumption S1 and Assumption S2(a)–(c), all initial penalized losses attain finite minima with probability at least* $1 - \delta/3$ *when* L *is sufficiently large and* a_n *sufficiently small.*

Proof of Lemma S2. For an initial outcome loss, let P denote its empirical average, $a_i = I(A_i = a)$, $u_i = \phi_i^T v$, $S_0 = Pa_i$, and $\nu_i = m_r^{a,\star}(\boldsymbol{X}_i)$. Write $v = (v_0, v_-^T)^T$, $\xi = P\{a_i(Y_i - \nu_i)\phi_i\}$. The scalar coercivity bound in Proposition S1 gives

$$Q(v) \geq cP(a_i|u_i|) - CS_0 - |\xi_0||v_0| - \|\xi_-\|_\infty \|v_-\|_1 + \lambda\|v_-\|_1.$$

Since $|v_0| \leq S_0^{-1}P(a_i|u_i|) + C_\phi\|v_-\|_1$, the events $|\xi_0|/S_0 \leq c/2$ and $\|\xi_-\|_\infty + C_\phi|\xi_0| \leq \lambda/2$ imply

$$Q(v) \geq (c/2)P(a_i|u_i|) + (\lambda/2)\|v_-\|_1 - CS_0.$$

This is coercive in (v_0, v_-) when $S_0 > 0$. Conditional Bernstein concentration and the positive arm probability give these events at the stated penalty and sample size. The same argument applies to the initial target logistic likelihood with response $I(A = a)$, conditional mean $e_t^{a,\star}(\boldsymbol{X})$, and $a_i = 1$.

For an initial source joint-weight, set $m = P_t\phi$ and $a_i = I(A_i = a)$ at the source. The bounded ideal weights satisfy $\mathbb{E}_s(a_i q_s^{a,\star}(\boldsymbol{X})\phi) = \mathbb{E}_t\phi$. Normalize their empirical version by $\widetilde{q}_i = q_s^{a,\star}(\boldsymbol{X}_i)/(P_s a_i q_s^{a,\star}(\boldsymbol{X}))$, so $P_s a_i \widetilde{q}_i = m_0 = 1$. On the bounded-moment concentration event the normalization is bounded above and away from zero. The identity $P_s a_i \widetilde{q}_i = 1$ and the bound $|\phi_{i,k}| \leq C_\phi$ control the extra normalization error by a constant times the intercept imbalance, so

$$\|m_- - P_s(a_i\widetilde{q}_i\phi_{i,-})\|_\infty \leq \lambda/2.$$

Writing the initial loss as

$$P_s\{a_i(e^{-u_i} + \widetilde{q}_i u_i)\} + \{m_- - P_s(a_i\widetilde{q}_i\phi_{i,-})\}^T v_- + \lambda\|v_-\|_1$$

shows it is at least $cP_s(a_i|u_i|) + (\lambda/2)\|v_-\|_1 - C$, again coercive. Here the ideal weights are used only for the proof and are not supplied to the algorithm. Bounded-coordinate concentration controls the normalization and moment imbalance at order a_n. A finite union bound over all initial-fitting blocks, with a sufficiently large L, proves the stated probability.

Continuity and coercivity give finite minimizers; uniqueness of the empirical minimizer is unnecessary. □

The following lemma verifies the preceding rate bound for the source–target calibration losses, including two-population balance, conditional-tail concentration, and predictor localization. For fixed positive constants C_0, C_1, write the sample-size conditions as

$$C_0 s_{\text{eff}} \sqrt{t_\delta / n} \le \kappa_G / 2, \qquad C_1 s_{\text{eff}} \lambda < r_0. \tag{S9}$$

Lemma S3 (Calibration rates)**.** *Under Assumptions S1 to S3, consider the prescribed pooled losses and penalties with unpenalized intercepts. For sufficiently large fixed L, C_0, C_1, Eq. (S9) gives the following event with probability at least $1 - \delta$: the correct initial fit and both calibrated updates attain finite minima, their predictors are bounded by $2M$, and the oracle bounds in Lemma S1 hold with $\kappa = \kappa_*$. The final fits satisfy*

$$\|\widehat{v}_c - \bar{v}_c\|_2 \le C\sqrt{s_c}\lambda, \qquad \|\widehat{v}_c - \bar{v}_c\|_1 \le C s_c \lambda,$$
$$\|\widehat{v}_w - \bar{v}_w\|_2 \le C\sqrt{s_w + s_{c,\text{init}}}\lambda, \qquad \|\widehat{v}_w - \bar{v}_w\|_1 \le C(s_w + s_{c,\text{init}})\lambda.$$

The target calibrated fits obey the same bounds. For bounded d and $\lambda \asymp n^{-1/2}$, finite minimizers and predictor localization hold with probability tending to one, with the displayed errors in O_p form.

Proof of Lemma S3. Lemma S2 supplies finite initial fits. Only the correctly specified initial fit needs an estimation rate; clipping controls the other fit throughout the conditional concentration argument. We allocate the remaining failure probability to the score and Gram events.

We first control the empirical gradients and Gram matrices. At each population target used in this lemma, its predictor has magnitude at most M. Each empirical gradient coordinate is either bounded or is a bounded multiple of Y plus bounded terms. The conditional Bernstein bound in Eq. (S4) and the bounded conditional mean therefore give a

uniform scalar Bernstein inequality. Apply this bound separately on each calibration block, conditional only on that block's own initial-training data; its clipped plug-in is then fixed and bounded. A union bound over d coordinates and F blocks, followed by summation with the prescribed sitewise weights, gives $\|A_N\|_\infty \le \lambda/2$ for L sufficiently large, except on an event of probability at most $\delta/3$. Initial correct-model gradients have zero population mean.

For each initial or calibration block, the entries of $I(A=a)\phi\phi^T$ are bounded by C_ϕ^2. Hoeffding concentration and a union bound over Fd^2 entries yield, with failure probability at most $\delta/3$,

$$\max_{r,a,\text{block}} \|\widehat{G}_{r,a} - G_{r,a}\|_{\max} \le CC_\phi^2\sqrt{t_\delta/n} =: e_n. \tag{S10}$$

Neither event assumes independence of the overlapping training sets.

These concentration events also give local curvature. For a fitted coefficient v within $\|v-\bar{v}\|_1 \le r_0$, every predictor along the connecting segment has magnitude at most $2M$. The source outcome Hessian is an empirical average of $I(A=a)e^{-\mathcal{T}_M(u_{\widehat{\gamma}_{\text{init}}})}g'(u_\alpha)\phi\phi^T$; the source joint-weight Hessian uses $I(A=a)e^{-u_\gamma}g'\{\mathcal{T}_M(u_{\widehat{\alpha}_{\text{init}}})\}\phi\phi^T$. The target calibrated Hessians have the same respective forms with β in place of γ. Initial outcome and joint-weight losses omit the extra plug-in factor; the initial target logistic likelihood uses $\sigma'(u_\beta)\phi\phi^T$. Each Hessian therefore dominates h_* times its empirical arm Gram matrix (or its full-target Gram matrix, which dominates an arm Gram matrix).

For any perturbation, Eq. (S10) gives

$$\Delta^T\widehat{G}_{r,a}\Delta \ge \kappa_G\|\Delta\|_2^2 - e_n\|\Delta\|_1^2.$$

The basic inequality implies $\|\Delta\|_1 \le (4\sqrt{s_v} + 2b_N/\lambda)\|\Delta\|_2$. For a correct initial fit $b_N = 0$. Once its rate is established, the propagation bound gives $b_N \le C\sqrt{s_{c,\text{init}}}\lambda$ for the opposite update, and $b_N \le C/N$ for a correct source joint-weight update, with zero conditional gradient for the other correct updates. Since $N^{-1}/\lambda = O(N^{-1/2})$, these cones all satisfy $\|\Delta\|_1 \le C\sqrt{s_{\text{eff}}}\|\Delta\|_2$. The first condition in Eq. (S9) makes their empirical Gram curvature

at least $\kappa_G/2$. Integrating the Hessian along the segment yields the loss remainder at least $h_*\kappa_G\|\Delta\|_2^2/4 = \kappa_*\|\Delta\|_2^2/2$.

The curvature bound now allows us to localize the fits. Consider first the correct initial fit, then the two calibrated fits in that order. Let $Q(v) = L_N(v) + \lambda\|v\|_{1,-0}$. If a point on the sphere $\|v - \bar{v}\|_1 = r_0$ had $Q(v) \le Q(\bar{v})$, its basic inequality would imply the applicable cone. The local curvature and the basic-inequality calculation in Lemma S1, applied to this boundary point, would then give $\|v - \bar{v}\|_1 \le C_1 s_{\text{eff}}\lambda < r_0$, a contradiction. Continuity ensures a minimizer over the closed ℓ_1 ball, and the boundary argument puts it in the interior. Convexity makes it a global minimizer. Any global minimizer outside the ball would, by convexity along its segment from $\bar{v}$, produce a boundary point with $Q(v) \le Q(\bar{v})$, again a contradiction. Thus every minimizer is localized. Applying the deterministic lemma inside the ball gives the stated rates. The correct initial rate supplies the propagated bound before it is used for the opposite update, so the argument is not circular.

For bounded d, the preceding probability bound with a fixed multiplier is replaced by a finite-dimensional argument. The coercivity inequalities in Proposition S1, together with bounded d and the fixed design, mean, and overlap constants, place all initial population minimizers in a common compact set. Their predictors are consequently uniformly bounded, although the misspecified initial bounds need not equal M. On a fixed neighborhood of each minimizer, strict positivity and continuity of g' and the exponential and logistic multipliers give a positive Hessian lower bound. The empirical arm Gram matrices converge, and each initial gradient at its population minimizer is $O_p(n^{-1/2})$ by the second-moment bounds. For a perturbation satisfying the basic inequality, local curvature thus gives

$$\frac{\kappa}{2}\|\Delta\|_2^2 \le \{\|\nabla L_N(\bar{v})\|_2 + \lambda\sqrt{d}\}\|\Delta\|_2.$$

The convex local-ball argument first yields finite initial minimizers and $O_p(n^{-1/2})$ errors. For the calibrated updates, blockwise conditional variances give $A_N = O_p(n^{-1/2})$; the

correct initial rate and Eq. (S7) give the same order for B_N. Their clipped multipliers have uniform positive lower bounds, so the same curvature and local-ball argument applies at the calibrated population coefficients. Hence every calibrated minimizer has ℓ_2 and ℓ_1 errors $O_p(n^{-1/2})$. The population margin then gives predictor localization with probability tending to one. Here stochastic gradient control replaces the high-probability gradient event used for growing d. □

For growing d, taking $\delta = (2d)^{-c}$ with fixed $c > 0$ gives $\lambda \asymp \sqrt{\log(2d)/N}$. The bounded-dimensional choice $\lambda \asymp N^{-1/2}$ has the same order. In both regimes, $s_{\text{eff}} \log(2d) = o(\sqrt{N})$ ensures the sample-size conditions in Eq. (S9).

The lemma yields $\|\widehat{v} - \bar{v}\|_1 = o_p(1)$ and hence

$$\sup_{\boldsymbol{x}} |\phi(\boldsymbol{x})^T(\widehat{v} - \bar{v})| \le C_\phi \|\widehat{v} - \bar{v}\|_1 = o_p(1).$$

The population margin therefore makes final target propensity clipping inactive with probability tending to one, uniformly over the fixed folds.

S2.8 Second-Order Bias and Score Stability

Orthogonality and Eq. (S5) convert coefficient error into second-order population bias and L_2 score error.

Lemma S4 (Second-order bias and score stability)**.** *Under Assumption S2 and Lemma 2, for $\vartheta \in \mathcal{N}$,*

$$|\mu_s^a(\vartheta) - \mu_s^a(\bar{\vartheta})| \le (C/2)\|\vartheta - \bar{\vartheta}\|_2^2,$$

$$\sum_{r \in \{t,s\}} \mathbb{E}_r\{h_r(D;\vartheta) - h_r(D;\bar{\vartheta})\}^2 \le C\|\vartheta - \bar{\vartheta}\|_2^2.$$

Proof of Lemma S4. For $\Delta = \vartheta - \bar{\vartheta}$, Taylor's formula along the segment from $\bar{\vartheta}$ to ϑ gives

$$\mu_s^a(\vartheta) - \mu_s^a(\bar{\vartheta}) = \nabla \mu_s^a(\bar{\vartheta})^T \Delta + \int_0^1 (1-u)\Delta^T \nabla^2 \mu_s^a(\bar{\vartheta} + u\Delta)\Delta \, du.$$

The first term is zero and the Hessian bound controls the integral. The second conclusion follows from the dimension-free bound in Eq. (S5). A general nonlinear GLM remainder includes squared errors of each nuisance; it need not reduce to their product under misspecification. □

S2.9 Candidate Expansion and Variance Consistency

Lemma 4 (restated)**.** *For the fitting procedure and penalty levels in Section 4.1, under Assumptions 1 to 3, each $s_j \in \mathcal{J}^{a,*}$ satisfies*

$$\widehat{\mu}_{s_j}^a - \mu_t^a = n_t^{-1} \sum_{i \in \mathcal{D}_t} \zeta_{t,s_j,i}^a + n_{s_j}^{-1} \sum_{i \in \mathcal{D}_{s_j}} \xi_{t,s_j,i}^a + o_p(N^{-1/2}),$$

and its fitted site-score covariances are consistent.

Proof of Lemma 4. Fix a source $s = s_j \in \mathcal{J}^{a,*}$. The conditions of Assumptions 2 and 3 give the pairwise design, tail, overlap, and population-margin conditions above. With the prescribed penalties, the sparsity condition implies

$$s_{\text{eff}} \lambda_N = o(1), \qquad \sqrt{N}\, s_{\text{eff}} \lambda_N^2 = o(1). \tag{S11}$$

Lemma S3 therefore gives $\|\widehat{\vartheta}^{(-k)} - \bar{\vartheta}\|_2 = O_p(\sqrt{s_{\text{eff}}}\lambda_N)$ and predictor localization on every fold. By Lemmas 1 and 2, $\mu_s^a(\bar{\vartheta}) = \mu_t^a$ and the population gradient vanishes.

For each evaluation fold k, set $\Delta h_{r,k} = h_r(D; \widehat{\vartheta}^{(-k)}) - h_r(D; \bar{\vartheta})$. Add and subtract the two population means to write the remainder as $K_f^{-1} \sum_k (R_{1,k} + R_{2,k}) + O_p(N^{-1})$, where the last term accounts for balanced fold sizes and

$$R_{1,k} = \sum_{r \in \{t,s\}} (\widehat{\mathbb{E}}_{r,k} - \mathbb{E}_r)\Delta h_{r,k}, \qquad R_{2,k} = \mu_s^a(\widehat{\vartheta}^{(-k)}) - \mu_s^a(\bar{\vartheta}).$$

Conditional on the training data for fold k, the evaluation terms are independent and centered. Since both site fold sizes are proportional to N,

$$\mathrm{Var}(R_{1,k} \mid \mathcal{D}^{(-k)}) \leq \frac{C}{N} \|\widehat{\vartheta}^{(-k)} - \bar{\vartheta}\|_2^2.$$

Conditional Chebyshev's inequality and consistency yield $R_{1,k} = o_p(N^{-1/2})$. The folds' training samples overlap, but no independence across the fixed number of fold remainders is needed to sum these bounds. By the bias lemma, $|R_{2,k}| = O_p(s_{\mathrm{eff}}\lambda_N^2) = o_p(N^{-1/2})$. This proves the expansion. The leading target and source sums are independent, site-centered averages. Uniform $(2+\delta)$ moments and positive limiting variance give the Lyapunov central limit theorem for their sum. Thus $\sqrt{N}(\widehat{\mu}_s^a - \mu_t^a) \xrightarrow{d} \mathcal{N}(0, V_{t,s}^{a,*})$.

To establish covariance consistency, center the cross-fitted scores within their sites and form

$$\widehat{V}_{t,s}^a = \frac{N}{n_t}\widehat{\mathbb{E}}_t((\widehat{\zeta}_{t,s}^a)^2) + \frac{N}{n_s}\widehat{\mathbb{E}}_s((\widehat{\xi}_{t,s}^a)^2) \xrightarrow{p} V_{t,s}^{a,*}.$$

Conditionally on each training sample, L_2 stability and Markov's inequality give $\widehat{\mathbb{E}}_{r,k}(\Delta h_{r,k})^2 = o_p(1)$. Summing over finitely many folds gives the same conclusion sitewise. Cauchy–Schwarz then shows that empirical first and second moments of fitted scores differ by $o_p(1)$ from their population-limit-score counterparts. The ordinary law of large numbers applies to the latter using their uniform integrability. Centering and the limiting site fractions complete the proof. The same argument applied to products gives consistency of covariances between two score components evaluated on the same observations.

The numbers of sources, arms, and folds are fixed, so the expansion and covariance conclusions hold simultaneously for all valid candidate components. □

S2.10 Outcome Links and Continuous Responses

For continuous outcomes with $g(u) = u$, $g'(u) = 1$ and $\Psi(u) = u^2/2$. The joint-weight calibration loss becomes the initial exponential-balancing loss, and the outcome calibration loss is weighted least squares with an ℓ_1 penalty. No initial outcome coefficient is required for the target derivative moment, which is simply $\widehat{\mathbb{E}}_t\phi$. The target regression average is $\widehat{\mathbb{E}}_t(\phi)^T\widehat{\alpha}^a_{t,s_j}$. With $\widehat{\Sigma}_{\phi,t} = \widehat{\mathrm{Cov}}_t\{\phi(\boldsymbol{X})\}$, the covariance between two source-specific target predictions is $(\widehat{\alpha}^a_{t,s_j})^T\widehat{\Sigma}_{\phi,t}\widehat{\alpha}^a_{t,s_h}$. Thus the same covariance structure applies, with particularly simple moment summaries. At the target, the propensity calibration also drops its initial-outcome multiplier, and the outcome update is weighted least squares with propensity-odds weights. The following corollary extends the inference result to unbounded continuous responses under conditional tail conditions.

Corollary S1 (Continuous-outcome inference)**.** *Take $g(u) = u$ and $\Psi(u) = u^2/2$. Under the design, conditional-tail, overlap, population-margin, and sparsity conditions of Lemma S3, together with Eq. (S11), the source-assisted mean has the expansion and consistent variance in Lemma 4 if either the source outcome mean is a correct linear model or the joint exponential weight is correct, with conditional-mean transportability in either case. At the target, impose the corresponding rate conditions with n_t and $s_{\mathrm{eff},0,a}$; either a correct linear outcome mean or a correct logistic propensity then yields the expansion and consistent variance in Proposition S2. These conclusions allow unbounded continuous Y. They also yield valid Wald inference for the treatment contrast, retaining cross-arm covariances.*

Proof of Corollary S1. The identity link has $g' = 1$ and $g'' = 0$, satisfying the link conditions on every bounded predictor interval. Source joint-weight calibration is the ordinary exponential-balancing loss and target propensity calibration has multiplier one. Outcome calibration is weighted least squares, with the same clipped initial source weights or target propensity-odds weights. Conditional residual centering and the respective balance equa-

tions give the casewise derivatives exactly as in Lemma 2 and Proposition S2. No step uses $Y \in \{0, 1\}$. Conditional Bernstein tails establish gradient concentration, and conditional second moments establish score stability and variance consistency. The verified nuisance rates and second-order argument then prove the stated expansions and normal limits. □

In particular, bounded conditional means with conditionally Gaussian errors of uniformly bounded variance satisfy the continuous-outcome conditions, without requiring equal variances across sites or treatment arms. The same general theorem covers Bernoulli outcomes with a logistic working mean and Poisson outcomes with an exponential working mean under the stated mean bounds. A working GLM specifies the nuisance mean and convex loss; arbitrary continuous distributions, heavy tails, and restricted-domain links are not automatically covered by that notation.

S2.11 From Arm Means to the Treatment Contrast

Apply the preceding expansion for $a = 1$ and $a = 0$, using their respective calibrated coefficients. Joint moment bounds give a bivariate central limit theorem by the Cramér–Wold argument. The pairwise treatment contrast consequently has variance

$$\frac{\mathrm{Var}_t(\zeta_{t,s}^1 - \zeta_{t,s}^0)}{n_t} + \frac{\mathrm{Var}_s(\xi_{t,s}^1 - \xi_{t,s}^0)}{n_s}$$

on the unscaled estimator scale. This expression includes both cross-arm covariances. Independence of the two potential-outcome mean estimators is not assumed.

S3 Joint Aggregation and Inference

We combine the source expansions with a target-only reference to prove Theorem 1. The steps below establish the target expansion, joint covariance, and stability of potentially biased candidates before source selection and oracle equivalence.

S3.1 Target-Only Limit Theory

Use the score, calibrated population coefficients, and losses defined in Section S1.2. The target-only reference has the following expansion.

Proposition S2 (Model-doubly-robust target expansion)**.** *Assume target causal identification, either a correct outcome GLM or a correct logistic propensity model, and the target versions of Assumptions S1 to S3. Use Eq. (S8) with $\delta_N \to 0$ when $d \to \infty$, or $\lambda_N \asymp n_t^{-1/2}$ when d is bounded, as covered by Lemma S3. Require $s_{\mathrm{eff},0,a}\lambda_N = o(1)$ and $\sqrt{n_t} s_{\mathrm{eff},0,a}\lambda_N^2 = o(1)$. For growing d and $\delta_N = (2d)^{-c}$, $s_{\mathrm{eff},0,a}\log(2d) = o(\sqrt{n_t})$ suffices. With $\varphi_t^a = h_0(D;\bar{\alpha}_0,\bar{\beta}_0) - \mu_t^a$, the target estimator satisfies*

$$\widehat{\mu}_t^a - \mu_t^a = n_t^{-1}\sum_{i\in\mathcal{D}_t}\varphi_{t,i}^a + o_p(n_t^{-1/2}).$$

It is asymptotically normal at rate $\sqrt{n_t}$, and the empirical variance of its centered cross-fitted scores is consistent for $\mathrm{Var}_t(\varphi_t^a)$ when this limit is positive. The corresponding bivariate result holds for both arms and their contrast.

Proof of Proposition S2. The target doubly robust identity gives $\mu_0(\bar{\alpha}_0,\bar{\beta}_0) = \mu_t^a$. Differentiation of the untruncated smooth functional gives

$$\nabla_\alpha \mu_0(\alpha,\beta) = \mathbb{E}_t[\{(1-I(A=a)) - I(A=a)e^{-\phi^T\beta}\}g'(\phi^T\alpha)\phi],$$

$$\nabla_\beta \mu_0(\alpha,\beta) = -\mathbb{E}_t\{I(A=a)e^{-\phi^T\beta}(Y - m(\boldsymbol{X};\alpha))\phi\}.$$

If outcome is correct, any positive clipped initial propensity-odds weight preserves the outcome minimizer $\alpha^\star$, and residual centering makes the propensity gradient zero for any β.

The outcome gradient is the stationarity equation of the target propensity loss with the correct outcome derivative; its clipping is inactive at truth.

If propensity is correct, $\mathbb{E}_t\{(1 - I(A = a)) - I(A = a)e^{-\phi^T\beta^\star} \mid \boldsymbol{X}\} = 0$. This both identifies the calibrated propensity under any clipped initial outcome multiplier and makes the outcome gradient vanish at any α. Because the true initial propensity predictor lies within the margin, the target outcome loss has weight $I(A = a)e^{-\phi^T\beta^\star}$; its stationarity makes the propensity gradient vanish at $\bar{\alpha}_0$. Thus both derivatives are zero in either branch, without equality of the misspecified initial and calibrated coefficients.

The conditional-gradient calculation and the verified rates of Lemma S3 apply to both target losses. The local target Hessian and L_2 bounds follow from Eq. (S5). Population bias is therefore $O_p(s_{\text{eff},0,a}\lambda_N^2) = o_p(n_t^{-1/2})$. On the event of uniform predictor localization, final propensity clipping is inactive; its complement has probability tending to zero and does not affect the expansion in probability. Cross-fitting bounds the conditional empirical-process remainder exactly as in the source proof. The ordinary central limit theorem, the cross-fitted second-moment argument, and Cramér–Wold prove the remaining claims. □

The limit score is evaluated at $(\bar{\alpha}_0, \bar{\beta}_0)$, not at a hypothetical pair of true functions when one working model is misspecified. Only when both models are correct does it reduce to the usual efficient target AIPW score. Model-doubly-robust inference here requires the specified calibrated fitting procedure and regularity; the same conclusion is not claimed for ordinary predictive nuisance fits.

S3.2 Covariance Decomposition

The covariance decomposition in Lemma 3 accounts for shared target observations and dependence between treatment arms.

Lemma 3 (restated)**.** *Independent site samples give additive variances for the first-order site contributions. The joint treatment-effect variance is estimated by*

$$\widehat{V}_{\tau}^{(-k_1)}(\eta) = \sum_{r} \frac{\widehat{\mathrm{Var}}_{r,-k_1}\{\widehat{G}^1_{r,i}(\eta^1) - \widehat{G}^0_{r,i}(\eta^0)\}}{n_{r,-k_1}}.$$

Here each empirical variance centers contributions within the inner calibration folds and pools their second central moments within the site. Under the candidate linear expansions and fitted-score covariance consistency, this expression consistently estimates the treatment-effect variance on the valid candidate subspace.

Proof of Lemma 3. Fix a pair of weight vectors supported on the valid candidate arms. The source expansion in Lemma 4 and the target expansion in Proposition S2 express their treatment contrast as a sum of centered site contributions, up to a negligible remainder. Samples from different sites are independent, so the variance of this sum is the sum of the site variances divided by their respective sample sizes.

Within a site, taking the variance of the difference between the two arm contributions preserves their cross-arm covariance. At the target, all candidate predictions are evaluated on the same observations, giving the full target covariance block. At each source, only that source's residual scores contribute, giving the diagonal source blocks across candidates. Expanding these within-site variances yields exactly the block formula in Section S1.3 and the expression in Eq. (10).

Under the assumed fitted-score covariance consistency, every entry of these blocks converges to its population counterpart on the valid candidate subspace. There are finitely many entries, and the variance is quadratic in the weights, so the corresponding scaled

variance converges uniformly on each compact set of valid weights. This establishes the claimed variance consistency. □

S3.3 Verification of Training Moments

To justify the covariance estimator and aggregation penalties, we verify Assumption 4 on the reused training sample. The proof applies population orthogonality and nuisance rates to an empirical Taylor expansion, then controls fitted-score second moments.

Proposition S3 (Training means and score covariances)**.** *For the primary pooled calibration procedure, assume the conditions of Lemma 4 for valid source arms and of Proposition S2 for the target. Let $n = n_{-k_1}$, let λ_N denote their common penalty order, and let $s_{\max}$ be the largest relevant effective support size. If $s_{\max}\lambda_N = o(1)$ and $\sqrt{n}\, s_{\max}\lambda_N^2 = o(1)$, the conditions in Assumption 4 hold uniformly over the fixed outer folds. Consequently, valid training discrepancies are $O_p(n^{-1/2})$ and their fitted-score covariance blocks are consistent, including shared-target and cross-arm products.*

Proof of Proposition S3. Fix one outer training sample and a valid source arm, suppressing $(-k_1)$ on fitted coefficients. Define its empirical functional by $\widehat{\mu}^a_{s,\mathrm{tr}}(\vartheta) = \widehat{\mathbb{E}}_{t,-k_1} h_t(D;\vartheta) + \widehat{\mathbb{E}}_{s,-k_1} h_s(D;\vartheta)$, and put $\Delta = \widehat{\vartheta}^{(-k_1)} - \bar{\vartheta}$. The rates in Lemma S3 give $\|\Delta\|_1 = O_p(s_{\max}\lambda_N)$ and $\|\Delta\|_2 = O_p(\sqrt{s_{\max}}\lambda_N)$, with predictors bounded by $2M$ along the connecting segment. The population gradient at $\bar{\vartheta}$ is zero by Lemma 2. Each empirical gradient coordinate at this deterministic coefficient is a sum of bounded terms and bounded multiples of $I(A = a)Y$. The conditional-tail bound and a coordinate union bound therefore give

$$\|\nabla \widehat{\mu}^a_{s,\mathrm{tr}}(\bar{\vartheta})\|_\infty = O_p(\lambda_N).$$

The gradient is evaluated at the population coefficient, so the bound remains valid when Δ depends on the training records.

To control the empirical Taylor remainder, use the target Gram matrix and the source matrix weighted by $I(A = a)(1 + |Y|)$. Their population operator norms are bounded: condition on $(\boldsymbol{X}, A)$ and apply the conditional first-moment and design bounds. The conditional-tail bound also gives uniformly subexponential tails for $I(A = a)|Y|$, so their maximum entrywise deviations are $O_p(\lambda_N)$ by Bernstein concentration. Every local Hessian block of $\widehat{\mu}^a_{s,\mathrm{tr}}$ is dominated in quadratic form by these two matrices, up to a fixed constant; mixed blocks are bounded using $2|xy| \le x^2 + y^2$. Uniformly along the segment, this yields

$$|\Delta^T \nabla^2 \widehat{\mu}^a_{s,\mathrm{tr}}(\vartheta)\Delta| \le C\|\Delta\|_2^2 + O_p(\lambda_N)\|\Delta\|_1^2.$$

Taylor's formula consequently gives

$$\begin{aligned} |\widehat{\mu}^a_{s,\mathrm{tr}}(\widehat{\vartheta}) - \widehat{\mu}^a_{s,\mathrm{tr}}(\bar{\vartheta})| &= O_p\{\lambda_N\|\Delta\|_1 + \|\Delta\|_2^2 + \lambda_N\|\Delta\|_1^2\} \\ &= O_p(s_{\max}\lambda_N^2 + s_{\max}^2\lambda_N^3) = O_p(s_{\max}\lambda_N^2). \end{aligned}$$

The last equality uses $s_{\max}\lambda_N = o(1)$. For the target-only functional, write the inverse propensity as $1 + e^{-\phi^T\beta_0}$. Its population gradient vanishes by the stationarity argument in Proposition S2, and its Hessian is controlled by the target Gram matrix weighted by $1 + I(A = a)|Y|$. The same Taylor bound therefore applies. Final propensity clipping is inactive on the localization event. Since $\sqrt{n}\, s_{\max}\lambda_N^2 = o(1)$, these bounds prove part (a) of Assumption 4.

For part (b), local Lipschitz bounds and bounded features give, for each applicable raw site contribution,

$$|h_r(D; \widehat{\vartheta}) - h_r(D; \bar{\vartheta})|^2 \le C\|\Delta\|_1^2\{1 + I(A = a)Y^2\}.$$

The empirical arm outcome second moments are $O_p(1)$, and $\|\Delta\|_1 = o_p(1)$, so the training averages of these squared differences are $o_p(1)$. The analogous bound holds for the target-only score. Centering fitted contributions within calibration folds preserves this conclusion: Jensen's inequality controls their mean differences, and each fold's population-coefficient

mean converges to its site expectation. Cauchy–Schwarz and the laws of large numbers for population-coefficient scores then give covariance consistency, including all shared-target and cross-arm entries. Finally, the population-nuisance training means have expectation μ_t^a and variance $O(n^{-1})$. The replacement bound makes their fitted discrepancies $O_p(n^{-1/2})$. A finite union over folds and valid arms completes the proof. □

S3.4 Candidate Consistency under Stable Initial Fits

Selection also uses moments from potentially biased candidates. These require consistency and a positive discrepancy scale. The following result supplies them under stability of both initial fits.

For a fixed source s_j and arm a, write $\widehat{\mu}_j = \widehat{\mu}_{s_j}^a$ and define $\bar{\alpha}_{\text{init}}, \bar{\gamma}_{\text{init}}$ by their initial population losses; the one-round $\bar{\alpha}_{\text{init}}$ is the target initial outcome minimizer. Define $\bar{\gamma}$ by the calibrated population joint-weight loss with plug-in $\bar{\alpha}_{\text{init}}$ and $\bar{\alpha}$ by the calibrated population outcome loss with plug-in $\bar{\gamma}_{\text{init}}$. Set $\bar{\vartheta} = (\bar{\alpha}^T, \bar{\gamma}^T)^T$. The corresponding functions are $\bar{m}_{s_j}^a(\boldsymbol{x}) = m(\boldsymbol{x}; \bar{\alpha})$ and $\bar{q}_{s_j}^a(\boldsymbol{x}) = q_{s_j}^a(\boldsymbol{x}; \bar{\gamma})$. These definitions retain the stated clipping, whether either model is correct or not. Let s_{stab} be the sum of all four support sizes, counting intercepts, and write

$$\mu_{j,N}^* = \mathbb{E}_t \bar{m}_{s_j}^a(\boldsymbol{X}) + \mathbb{E}_{s_j}\{I(A = a)\bar{q}_{s_j}^a(\boldsymbol{X})(Y - \bar{m}_{s_j}^a(\boldsymbol{X}))\}.$$

Proposition S4 (Stable candidate means and score covariances)**.** *Suppose the population design, GLM, conditional-tail, and quantitative-overlap conditions in Assumptions S1 and S2(a)–(c) hold, without assuming transportability or either correct nuisance model. Assume all four initial and calibrated population predictors are bounded by M, and $s_{\text{stab}}\lambda_N = o(1)$ with the penalty choice in Eq. (S8), $\delta_N \to 0$, for growing d, or $\lambda_N \asymp N^{-1/2}$ for bounded d. The complete outer-evaluation estimate satisfies*

$$\widehat{\mu}_j - \mu_{j,N}^* = O_p\{\sqrt{s_{\text{stab}}}\lambda_N + N^{-1/2}\} = o_p(1).$$

On both training and evaluation samples, fitted sitewise score means, variances, and cross-products differ by $o_p(1)$ from their population-coefficient counterparts. The reused training candidate means satisfy $\widehat{\mu}_j - \mu^*_{j,N} = o_p(1)$. Cross-products require the corresponding conditions for the other candidate, using Proposition S2 for the target-only reference. When the population means and score covariances have limits, these empirical moments converge to them. If additionally $\mathrm{Var}_{s_j}(Y \mid \boldsymbol{X}, A = a) \geq v_- > 0$, the discrepancy standard error satisfies $\sqrt{n_{-k_1}}\, \widehat{s}^a_{s_j} \to_p s^{a,*}_{s_j} \in (0, \infty)$.

Proof of Proposition S4. The population targets exist uniquely by Proposition S1. At each initial population minimizer the gradient is zero, even under misspecification, and the predictor bound gives the same Bernstein concentration as in Lemma S3. Its Gram and local-ball arguments with $b_N = 0$ give initial ℓ_2 errors $O_p(\sqrt{s_{\mathrm{stab}}}\lambda_N)$ and ℓ_1 errors $O_p(s_{\mathrm{stab}}\lambda_N)$, since each initial support is included in s_{stab}.

For either calibrated update, compare its conditional population gradient with the gradient using its own limiting initial plug-in. With that common population plug-in, each site's fold weights sum to one, so pooling recovers the population calibration loss and its zero gradient. The Lipschitz propagation calculation in Lemma S1 therefore gives $b_N = O_p(\sqrt{s_{\mathrm{stab}}}\lambda_N)$ without requiring model correctness. The expanded cone has $\|\Delta\|_1 \leq C\sqrt{s_{\mathrm{stab}}}\|\Delta\|_2$. Since $s_{\mathrm{stab}}\lambda_N = o(1)$, the Gram and localization conditions hold eventually, and the same local-ball argument gives final ℓ_2 errors $O_p(\sqrt{s_{\mathrm{stab}}}\lambda_N)$ and ℓ_1 errors $O_p(s_{\mathrm{stab}}\lambda_N)$. For bounded d, the finite-dimensional argument in Lemma S3 applies at all four population minimizers and gives the same rates.

On the resulting local event, the population functional is Lipschitz in coefficient ℓ_2 norm. Each gradient is an expectation of ϕ times a scalar with bounded conditional mean;

Cauchy–Schwarz and the population Gram upper bound control every unit directional derivative. Thus, uniformly in the local neighborhood,

$$\|\nabla \mu^a_{s_j}(\vartheta)\|_2 \le C, \qquad |\mu^a_{s_j}(\widehat{\vartheta}) - \mu^*_{j,N}| \le C\|\widehat{\vartheta} - \bar{\vartheta}\|_2.$$

This population error is $O_p(\sqrt{s_{\text{stab}}}\lambda_N)$; orthogonality is not required. Conditional on each outer training sample, the held-out site averages have bounded variance of order N^{-1}. As in the proof of Lemma 4, replacing balanced sitewise fold fractions by $1/K_f$ costs $O_p(N^{-1})$. Summing the fixed number of fold terms gives the stated outer-evaluation rate.

On each outer evaluation fold, the L_2 score bound in Eq. (S5), conditional Markov's inequality, and Cauchy–Schwarz make empirical fitted-score first and second moments differ by $o_p(1)$ from their population-coefficient counterparts. The population-predictor bound supplies the local neighborhood used in that inequality. Uniform conditional outcome moments give the corresponding laws of large numbers. The same argument applies to products with another stable score, including the target-only score.

On reused training observations, the ℓ_1 rates and bounded features give uniform $o_p(1)$ predictor errors. Local Lipschitz bounds then control each squared score error by a common $o_p(1)$ factor times $1 + Y_i^2$. The sitewise empirical outcome second moments are $O_p(1)$, so the same moment consistency holds on these training samples without conditioning on the fitted model. The fixed number of internal training constructions preserves these bounds.

Finally, for $U_j = I(A = a)\bar{q}^a_{s_j}(\boldsymbol{X})(Y - \bar{m}^a_{s_j}(\boldsymbol{X}))$, the law of total variance gives

$$\mathrm{Var}_{s_j}(U_j) \ge \mathbb{E}_{s_j}\{I(A = a)\bar{q}^a_{s_j}(\boldsymbol{X})^2 \,\mathrm{Var}(Y \mid \boldsymbol{X}, A = a)\} \ge e^{-2M} v_- P_{s_j}(A = a).$$

Because the intercept is in the design, Eq. (S3) implies $P_{s_j}(A = a) \ge \kappa_G$. Under the stated covariance limits, the limit of $n_{-k_1}(\widehat{s}^a_{s_j})^2$ is the source variance divided by $\rho^{\text{all}}_{s_j}$, plus a nonnegative target-score variance. The uniform second-moment bounds make this limit finite. Balanced folds preserve these site fractions, so

$$(s^{a,*}_{s_j})^2 \ge \frac{e^{-2M} v_- \kappa_G}{\rho^{\text{all}}_{s_j}} > 0. \qquad \square$$

The fixed-active-law construction in Proposition S1 supplies stable sparse targets even when site outcome means differ. We now combine this candidate stability with the valid-candidate training moments to prove selection and oracle equivalence.

S3.5 Proof of the Joint Aggregation Theorem

Work on an outer training sample of size $n = n_{-k_1}$, suppressing $(-k_1)$ until evaluation. All statements hold uniformly over the fixed number of outer folds.

Theorem 1 (restated)**.** *For the fitting procedure and penalty levels in Section 4.1, under Assumptions 1 to 5 and the target-only conditions of Proposition S2, the outer-fold weights satisfy* $\widehat{\eta}^{(-k_1)} \to_p \eta^*$ *and*

$$P\{\widehat{\eta}_{s_j}^{a,(-k_1)} = 0 \text{ for all } a = 0, 1,\ s_j \notin \mathcal{J}^{a,*}\} \to 1.$$

Moreover,

$$\sqrt{N}(\widehat{\tau}_t - \tau_t) \xrightarrow{d} \mathcal{N}(0, V_\tau^*), \qquad V_\tau^* \le L_\tau^*(0).$$

If $\mathcal{J}^{0,*} = \mathcal{J}^{1,*} = \varnothing$*, the aggregate equals the target-only contrast with probability tending to one.*

Proof of Theorem 1. The quadratic representation in Assumption 5 has $\widehat{Q}_n \to_p Q^*$ with $\lambda_{\min}(Q^*) > 0$ and $\widehat{\boldsymbol{l}}_n = O_p(1)$. On an event of probability tending to one, $\lambda_{\min}(\widehat{Q}_n) \ge \kappa$ for some fixed $\kappa > 0$. Comparing the penalized objective at $\widehat{\eta}$ with its value at zero, and using nonnegative penalty coefficients, gives

$$\kappa\|\widehat{\eta}\|_2^2 \le 2\|\widehat{\boldsymbol{l}}_n\|_2\|\widehat{\eta}\|_2, \qquad \|\widehat{\eta}\|_2 = O_p(1).$$

This argument controls the unconstrained optimizer without imposing a compact weight set.

We next examine the penalty on each coordinate. For a valid source arm, Proposition S3 gives a training-mean nuisance replacement error $o_p(n^{-1/2})$. The centered population contributions have finite second moments by Assumption 2, and their expectation equals μ_t^a. Hence $\widehat{d}_{s_j}^a = O_p(n^{-1/2})$ and $\widehat{t}_{s_j}^a = O_p(1)$. Its penalty coefficient $(\widehat{t}_{s_j}^a/c_n - 1)_+$ is therefore zero with probability tending to one.

For a biased arm, $\widehat{d}_{s_j}^a \to_p \bar{d}_{s_j}^a \neq 0$ and $\sqrt{n}\widehat{s}_{s_j}^a \to_p s_{s_j}^{a,*} > 0$. Consequently $\widehat{t}_{s_j}^a = \Theta_p(\sqrt{n})$ and, since $c_n/\sqrt{n} \to 0$, its penalty coefficient diverges. The quadratic gradient at the optimizer is $2\widehat{\boldsymbol{l}}_n + 2\widehat{Q}_n\widehat{\eta} = O_p(1)$. A nonzero biased coordinate would require its gradient component to equal the diverging penalty coefficient in absolute value, contradicting this stochastic bound. There are finitely many source arms, so all biased coordinates are zero with probability tending to one. On the remaining valid coordinates the penalty is zero, and quadratic coefficient convergence with positive curvature gives the unique restricted minimizer $\widehat{\eta} \to_p \eta^*$.

Restore the outer fold index k_1. Form a signed discrepancy vector from held-out valid source-minus-target means, with negative signs in the control-arm coordinates and zeros on biased coordinates. The calibrated source expansion in Lemma 4 and the target-anchor expansion in Proposition S2 give $\widehat{\boldsymbol{D}}_{k_1} = O_p(N^{-1/2})$. The foldwise difference between the fitted-weight and oracle-weight contrasts is therefore bounded by

$$|(\widehat{\eta}^{(-k_1)} - \eta^*)^T \widehat{\boldsymbol{D}}_{k_1}| \leq \|\widehat{\eta}^{(-k_1)} - \eta^*\|_2 \|\widehat{\boldsymbol{D}}_{k_1}\|_2 = o_p(N^{-1/2}).$$

The balanced folds satisfy $n_{r,k_1}/n_r = K_f^{-1} + O(n_r^{-1})$. Their site fractions therefore differ by $O(N^{-1})$, which contributes $o_p(N^{-1/2})$ to the pooled fitted-weight difference. The finite number of folds preserves the preceding product order. Independence between the learned weights and held-out candidate differences is unnecessary for this bound.

At the oracle weights, the candidate expansions yield site-centered population contributions G_r^a as defined in the main text. Pooling folds assigns each observation coefficient $1/n_r$ at site r. Site independence and the $(2+\delta)$ moment bounds give

$$\sqrt{N}(\widehat{\tau}_t - \tau_t) \xrightarrow{d} \mathcal{N}(0, V_\tau^*), \qquad V_\tau^* = \sum_r \frac{\mathrm{Var}_r(G_r^1 - G_r^0)}{\rho_r^{\mathrm{all}}}.$$

The fitted-score covariance argument in the proof of Lemma 4 and the target score consistency in Proposition S2, together with $\widehat{\eta} \to_p \eta^*$, imply empirical L_2 convergence of the outer contributions to these population contributions. Their pooled second moments then give $N\widehat{V}_\tau^{\mathrm{out}} \to_p V_\tau^*$. Finally, zero weights are feasible for the oracle objective, so $V_\tau^* \le L_\tau^*(0)$. If both valid source sets are empty, selection sets every source weight to zero with probability tending to one, leaving the target-only contrast. □

S3.6 Reported Standard Errors

The preceding proof establishes consistency of the leading variance in Eq. (13). The simulation and RHC analyses use its square root as the primary RoCE standard error, with the learned weights plugged into the site-centered outer contributions. Their Wald intervals do not include an additional weight-learning correction. The All of Us figure retains its supplied intervals pending a matching export.

Optional local sensitivity diagnostic. The implementation also retains a local delta-method calculation for estimating the aggregation weights. It is reported separately from the primary standard error. We derive it by perturbing training moments and differentiating the weight optimality equations.

The calculation holds nuisance fits, feature transformations, folds, tuning values, and the selected active set fixed. It describes the reported local sensitivity calculation; inference at selection boundaries or under full refitting requires additional analysis.

Fix outer fold k and write $n_{r,-k}$ for its training count at site r and $n_{-k} = \sum_r n_{r,-k}$. Stack the uncentered training records as $W_{r,i,k} = (B_{r,i,k}, U_{r,i,k}^T)^T$, with $U_{r,i,k} \in \mathbb{R}^{2K}$. At the target, $B_{t,i,k} = \widehat{h}^1_{0,i} - \widehat{h}^0_{0,i}$ and

$$U_{t,i,k} = (\widehat{h}^1_{t,1,i} - \widehat{h}^1_{0,i}, \ldots, \widehat{h}^1_{t,K,i} - \widehat{h}^1_{0,i}, -\widehat{h}^0_{t,1,i} + \widehat{h}^0_{0,i}, \ldots, -\widehat{h}^0_{t,K,i} + \widehat{h}^0_{0,i})^T.$$

At source s_j, $B_{s_j,i,k} = 0$; $U_{s_j,i,k}$ has only two potentially nonzero entries: $\widehat{h}^1_{s_j,j,i}$ in coordinate j and $-\widehat{h}^0_{s_j,j,i}$ in coordinate $K + j$. All fitted functions exclude outer fold k. These are the scores of Eq. (S2), expressed as a target-only contrast plus signed source increments.

Let $\bar{W}_{r,k}$ be the pooled site mean and $\widetilde{W}_{r,i,k}$ the record centered within its inner calibration fold. Define

$$\widehat{C}_{r,k} = \frac{1}{n_{r,-k}} \sum_{i \in \mathcal{I}_{r,-k}} \widetilde{W}_{r,i,k} \widetilde{W}_{r,i,k}^T, \qquad \begin{pmatrix} \widehat{v}_k & \widehat{l}_k^T \\ \widehat{l}_k & \widehat{H}_k \end{pmatrix} = \sum_r \frac{\widehat{C}_{r,k}}{n_{r,-k}}.$$

Thus $\widehat{v}_k + 2\widehat{l}_k^T \eta + \eta^T \widehat{H}_k \eta$ is the unscaled variance objective. In the main theorem's scaled notation, $\widehat{Q}_n = n_{-k} \widehat{H}_k$. The signed discrepancy vector is $\widehat{\delta}_k = \sum_r \bar{U}_{r,k}$. The numerical implementation uses the variance floor $\epsilon_V = 10^{-6}$. Set

$$\widehat{s}_{k,j} = \{\max(\widehat{H}_{k,jj}, \epsilon_V)\}^{1/2}, \quad \widehat{t}_{k,j} = |\widehat{\delta}_{k,j}| / \widehat{s}_{k,j}, \quad p_{k,j} = (\widehat{t}_{k,j}/c - 1)_+.$$

The reported analyses use $c = 2$. The sign convention for control-arm discrepancies does not affect their absolute Wald statistics.

To differentiate the moments, perturb the empirical mass of one training observation i at site r from one to $1 + \varepsilon$, retaining the original sample counts in the variance scaling. Means and within-fold centering use the perturbed masses. A dot denotes differentiation at $\varepsilon = 0$. For the covariance blocks $\widehat{C}^{UU}_{r,k}$ and $\widehat{C}^{UB}_{r,k}$, the moment derivatives satisfy

$$n_{r,-k} \dot{\delta}_k = U_{r,i,k} - \bar{U}_{r,k},$$

$$n_{r,-k}^2 \dot{H}_k = \widetilde{U}_{r,i,k} \widetilde{U}_{r,i,k}^T - \widehat{C}^{UU}_{r,k},$$

$$n_{r,-k}^2 \dot{l}_k = \widetilde{U}_{r,i,k} \widetilde{B}_{r,i,k} - \widehat{C}^{UB}_{r,k}.$$

Only coordinates with $\widehat{t}_{k,j} > c$ contribute a penalty derivative. Away from the variance-floor kink, the chain rule gives

$$\dot{s}_{k,j} = \frac{I(\widehat{H}_{k,jj} > \epsilon_V)\dot{H}_{k,jj}}{2\widehat{s}_{k,j}}, \qquad \dot{t}_{k,j} = \frac{\operatorname{sign}(\widehat{\delta}_{k,j})\dot{\delta}_{k,j} - \widehat{t}_{k,j}\dot{s}_{k,j}}{\widehat{s}_{k,j}}.$$

Hence $\dot{p}_{k,j} = \dot{t}_{k,j}/c$ in this region, whereas $\dot{p}_{k,j} = 0$ when $\widehat{t}_{k,j} < c$. Below the variance floor, $\dot{s}_{k,j} = 0$.

Let J_k be the selected nonzero coordinates, with their signs held fixed. Numerically, coordinates with $|\widehat{\eta}_{k,j}| > 10^{-12}$ are treated as active, and $|\widehat{t}_{k,j}/c - 1| < 10^{-6}$ flags proximity to a penalty kink. Differentiating the stationarity equations of the soft-penalized objective yields

$$\widehat{H}_{k,J_kJ_k}\dot{\eta}_{k,J_k} = -(\dot{H}_k\widehat{\eta}_k)_{J_k} - \dot{l}_{k,J_k} - (2n_{-k})^{-1}\operatorname{diag}\{\operatorname{sign}(\widehat{\eta}_{k,J_k})\}\dot{p}_{k,J_k}.$$

Inactive coordinates have $\dot{\eta}_{k,J_k^c} = 0$; for an empty active set, the entire derivative is zero. The calculation requires nonsingular active curvature and does not differentiate a change in the selected coordinates.

The effect of each weight perturbation on the final estimate depends on the actual outer-fold increments. For $a \in \{0, 1\}$ define

$$v_{k,j}^a = \frac{1}{n_t}\sum_{i\in\mathcal{I}_{t,k}}(\widehat{h}_{t,j,i}^a - \widehat{h}_{0,i}^a) + \frac{1}{n_{s_j}}\sum_{i\in\mathcal{I}_{s_j,k}}\widehat{h}_{s_j,j,i}^a.$$

Stack these as $v_k = ((v_k^1)^T, -(v_k^0)^T)^T$. The functions are evaluated on outer fold k, and the denominators are the full site counts, as in Eq. (12). Accumulate an observation's indirect contribution over every outer training set containing it:

$$b_{r,i} = \sum_{k:\, i\in\mathcal{I}_{r,-k}} v_k^T\dot{\eta}_{k;r,i}.$$

Let $T_{r,i} = \widehat{G}_{r,i}^1 - \widehat{G}_{r,i}^0$ be its held-out contrast contribution and set $a_{r,i} = (T_{r,i} - \bar{T}_r)/n_r$. The leading variance and the separate local diagnostic variance are

$$\widehat{V}_\tau^{\text{out}} = \sum_r\sum_{i\in\mathcal{D}_r} a_{r,i}^2, \qquad \widehat{V}_{\text{diag}} = \sum_r\sum_{i\in\mathcal{D}_r}(a_{r,i} + b_{r,i})^2. \tag{S12}$$

Using empirical second moments with divisor n_r, $\sum_{r,i} a_{r,i}^2$ equals the leading fixed-weight variance. The diagnostic adds both $\sum_{r,i} b_{r,i}^2$ and $2\sum_{r,i} a_{r,i} b_{r,i}$, and hence need not increase the standard error. Primary intervals use $\widehat{V}_\tau^{\text{out}}$, not $\widehat{V}_{\text{diag}}$. Accumulating by original observation before squaring retains dependence across overlapping training folds. This local calculation neither refits nuisance models nor accounts for discrete selection changes; it does not resolve the weak-bias coverage limitation discussed in the main paper.

The first- and second-moment messages in figure 1 suffice for point estimation and the leading variance. The optional diagnostic above uses observation-level score records in the numerical implementation. Computing it through federated summaries would also require the higher-order and cross-fold products entering Eq. (S12); those summaries are beyond the message list in that figure.

S4 Simulation Models and Comparators

This section supplies the model and comparator details for Section 5. The simulations vary nuisance misspecification and source incompatibility separately. Every fitted model uses an intercept and the $p = 200$ covariate main effects. The generated interaction X_1X_2 enters through $z_\kappa = X_1(1 + \kappa X_2)/\sqrt{1 + \kappa^2}$ whenever the corresponding κ is nonzero. The outcome interaction is controlled by κ_Y, while κ_A controls interactions in treatment assignment and population shift. The statuses in table S1 concern $\rho = 0$; for $\rho > 0$, source 1's treated outcome additionally violates transportability. A correct outcome functional form at a site does not remove that violation.

Table S1: Generated and fitted nuisance models in the three simulation scenarios at $\rho = 0$. All fitted models omit X_1X_2.

Scenario	κ_Y	κ_A	Outcome model	Treatment and joint-weight models
1	0	0	Correct	Correct
2	0.75	0	Misspecified: omitted interaction	Correct
3	0	0.75	Correct	Misspecified: omitted interaction

The benchmark names refer to the implementations compared in this manuscript. They combine familiar AIPW and federated aggregation ideas (Bang and Robins, 2005; Xiong et al., 2023), with the following distinctions:

Calibrated target-only.

Calibrated target outcome and propensity models produce an AIPW treatment contrast using target observations alone. This is the reference for assessing the gain from source borrowing.

Sample-size aggregation (SS).

Local arm-mean estimates are combined according to site sample sizes. The procedure does not transport the source arm means to the target covariate distribution.

Inverse-variance aggregation (IVW).

Local arm-mean estimates are combined using inverse-variance weights, including a nonnegative between-site variance adjustment. As with SS, the source means are not transported to the target.

Federated-DR.

Density-ratio-weighted source AIPW estimates are combined with the target estimate

by inverse variance. Treatment and population adjustments use separate nuisance components, and there is no RoCE pivotal penalty.

Pooled-DR.

Density-ratio-weighted AIPW uses common nuisance fits obtained from pooled records. This comparator provides a pooled-data reference and does not satisfy the same record-sharing restriction as the federated procedure.

RoCE combines paired calibration with arm-specific pivotal penalties and joint TATE variance minimization. The comparisons evaluate these complete procedures; differences cannot be attributed to a single component in the absence of a corresponding ablation. All methods use the same generated dataset in each replicate. RMSE measures the combined effect of bias and variability, while coverage assesses the reported intervals. Pointwise Monte Carlo intervals describe simulation uncertainty and are not simultaneous bands across settings.

S5 All of Us Cohort Definition and Site Composition

This section supplements Section 6 with the cohort rules and site composition used in the All of Us analysis.

The analysis defines an initial-exposure contrast in the evaluable one-year cohort. Index is the earlier first recorded exposure to GLP-1 or DPP-4; simultaneous first exposures are assigned to DPP-4. Participants must have a type 2 diabetes diagnosis by index and no prior diagnosis in the composite outcome set. Subsequent exposure to the other class does not change the initial assignment. The design imposes no separate washout or background-therapy restriction.

[To be completed: list the diseases or diagnostic categories in the composite dermatologic outcome, the corresponding diagnosis codes or concept identifiers, and the rule used to combine them.]

The composite outcome included any of the nine dermatologic conditions listed in table S2. These conditions were identified using their SNOMED CT concepts, represented by the OMOP concept IDs in table S2, together with all descendant concepts in the All of Us condition-concept hierarchy.

Table S2: Starting concepts for the composite dermatologic outcome. Each concept and its descendants are included.

Disease or diagnostic category	OMOP concept ID
Hidradenitis suppurativa	4241223
Inflammatory dermatosis	45766714
Psoriasis	140168
Cutaneous lupus erythematosus	4324123
Eczema	133835
Urticaria	139900
Bullous pemphigoid	4298692
Pemphigus vulgaris	4170723
Alopecia areata	141933

For each participant, the earliest recorded diagnosis matching any of these concepts or their descendants defined composite-outcome onset, with participants whose onset preceded index excluded from the cohort. Among those retained, a diagnosis of any component condition between index and one calendar year later, including both endpoints, constituted a composite event. For participants without such an event, a negative outcome was assigned only when the last condition-occurrence date was on or after the one-year endpoint.

Of 15,339 baseline participants with geographic information, 2,074 have unclassifiable outcomes and are excluded, leaving 13,265 participants. No censoring weights are used. This rule defines the evaluable population; it does not establish continuous observation throughout the year.

Covariates comprise age, ethnicity, gender, race, sex at birth, and diagnosis-derived PheCode counts for the year before index and earlier history. Diabetes and outcome-defining concepts are excluded. The 5% prevalence screen precedes splitting and yields 154 count features among 181 model features. Scaling uses outer training samples. Northeast is the largest evaluable geographic group and is designated as the target across all methods. Table S3 describes the four groups; the marginal event counts are descriptive and are not tests of conditional transportability.

Table S3: All of Us evaluable one-year cohort by geographic site. Treatment-arm entries give participants and recorded composite-outcome events in parentheses. These unadjusted counts are distinct from the adjusted risk differences in figure 4.

Site	Region	Total n	GLP-1 n (events)	DPP-4 n (events)
t	Northeast	5,040	1,082 (90)	3,958 (277)
s_1	Midwest	3,433	874 (84)	2,559 (205)
s_2	South	2,429	567 (42)	1,862 (136)
s_3	West	2,363	588 (39)	1,775 (118)

S6 Right Heart Catheterization Analysis

This second application complements the All of Us analysis in Section 6 using an insurance-defined target population.

We analyze 30-day mortality after right heart catheterization (RHC) in 5,039 SUPPORT patients (Connors et al., 1996; Hirano and Imbens, 2001). Private insurance defines the target group ($n_t = 1{,}698$); Medicare, private insurance with Medicare, and Medicaid define three candidate sources. Uninsured patients and those with both Medicare and Medicaid are excluded. All methods use the same cohort and 61 covariates plus an intercept. RoCE and calibrated target-only use two-layer ten-fold cross-fitting; median/mode imputation and standardization are fitted on each pooled outer training sample. After this pooled preprocessing, the implemented joint-TATE analysis uses one batched source-model exchange, a fixed cutoff of 2, and numerical predictor radii of 3 at sources and 5 at the target. Pooled medians and modes may require additional communication in a genuinely distributed implementation. SS and IVW use site-specific training-fold preprocessing, while the two DR benchmarks use full-sample fitting. These application-specific radii differ from the common-radius sufficient regime in the main paper.

Figure S1 shows six mortality risk-difference estimates. RoCE gives 3.02 percentage points (95% CI $[-1.13, 7.16]$), versus 2.40 points (95% CI $[-2.42, 7.22]$) for calibrated target-only estimation. Their standard errors are 2.12 and 2.46 points, respectively, for an estimated relative efficiency of 1.35. The RoCE interval uses the leading variance; reaggregation of the saved fit leaves the point estimate and source weights unchanged and requires no nuisance-model refit. The insurance groups emulate sites within one study, not independently recruited hospitals. Target-anchor weights averaged across outer folds are approximately 0.671 for the treated mean and 0.729 for the control mean (table S4).

Table S4: RHC insurance groups and RoCE source weights averaged over ten outer folds. RHC and mortality are observed proportions.

Insurance	n	RHC	Mortality	$\widehat{\eta}^1_{s_j}$	$\widehat{\eta}^0_{s_j}$
Private (target)	1698	0.431	0.309	–	–
Medicare	1458	0.350	0.362	0.092	0.111
Private & Medicare	1236	0.396	0.367	0.087	0.067
Medicaid	647	0.298	0.269	0.149	0.093

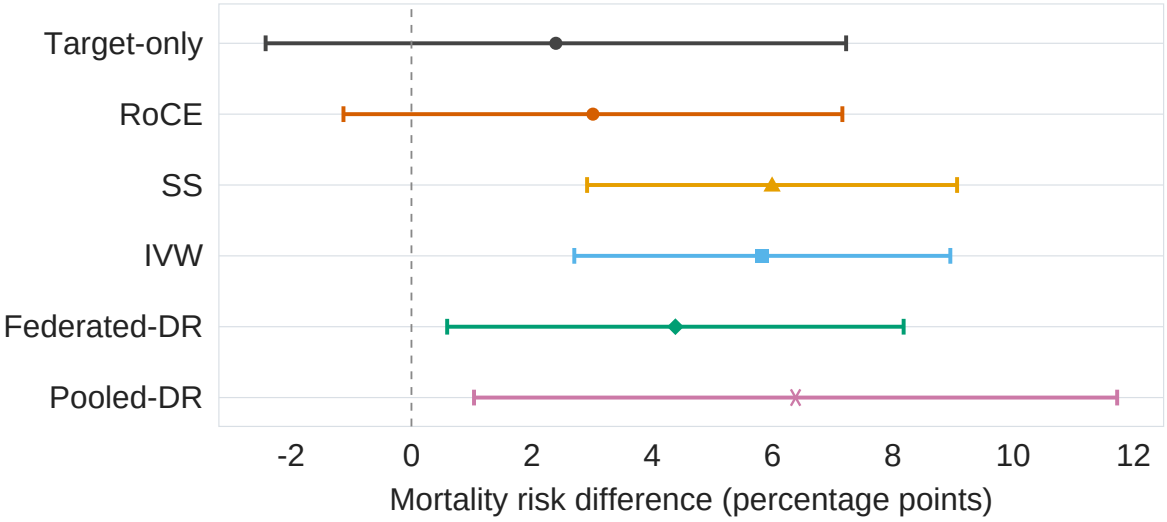


Figure S1: Estimated target-population effects of RHC on 30-day mortality with 95% confidence intervals. RoCE uses the leading site-centered variance; baseline intervals retain their original inference procedures. Positive risk differences indicate higher mortality under RHC; the dashed line marks zero.

S7 Alternative Validation and Armwise Aggregation

We consider two alternatives to the primary procedure: independent inner validation for weight learning, and armwise variance minimization on a bounded weight set. Unless both arms are displayed, fix a, write $\mu = \mu_t^a$, and suppress the arm index.

S7.1 Independent Inner Validation

Partition each site into $K_f \geq 4$ balanced folds $\mathcal{I}_{r,h}$. Reserve k for outer evaluation and $\ell \neq k$ for inner validation. The remaining index h denotes a calibration fold. For an excluded set B, define $\mathcal{T}_r(B) = \bigcup_{h \notin B} \mathcal{I}_{r,h}$ and $n_r(B) = |\mathcal{T}_r(B)|$. Outer refits use $B = \{k\}$; inner validation fits use $B = \{k, \ell\}$. Initial fits for block $h \notin B$ use only $\mathcal{T}_r(B \cup \{h\})$. The same exclusion applies to target summaries, preprocessing, and tuning.

Write $L_{\gamma,h}$ and $L_{\alpha,h}$ for the unpenalized empirical loss contributions in Eq. (4) and Eq. (5), evaluated on calibration fold h. The complete calibration fit is

$$\begin{aligned} \widehat{\gamma}^{(-B)} &\in \arg\min_{\gamma} \left\{ \frac{1}{K_f - |B|} \sum_{h \notin B} L_{\gamma,h}(\gamma; \widehat{\alpha}_{\text{init}}^{(-B,-h)}) + \lambda_\gamma(B) \|\gamma\|_{1,-0} \right\}, \\ \widehat{\alpha}^{(-B)} &\in \arg\min_{\alpha} \left\{ \frac{1}{K_f - |B|} \sum_{h \notin B} L_{\alpha,h}(\alpha; \widehat{\gamma}_{\text{init}}^{(-B,-h)}) + \lambda_\alpha(B) \|\alpha\|_{1,-0} \right\}. \end{aligned} \tag{S13}$$

Each site average uses all observations in its calibration fold; $I(A = a)$ selects the treatment arm without changing the denominator. For unequal folds, weight each site's contribution by its own fraction $|\mathcal{I}_{r,h}|/n_r(B)$ before combining the target and source terms. Add each penalty once. The target derivative moment uses target fold h and the initial outcome fit excluding $B \cup \{h\}$; the feature mean for an initial joint-weight fit uses only the initial-training target sample.

Apply the same construction to the target losses in Eq. (S1), stated with their primary fold updates in Section S1.2. Initial target outcome fits use the corresponding treatment arm;

initial logistic propensity fits use all target training observations with response $I(A = a)$. Calibrated target propensities are $\sigma_M(\phi^T \widehat{\beta}_{0,a})$.

For inner validation fold ℓ, evaluate only the calibrated models with $B = \{k, \ell\}$. Denote the resulting uncentered target-only contribution by $h_{0,i}^{(-k,-\ell)}$, target prediction for source j by $h_{t,j,i}^{(-k,-\ell)}$, and source residual contribution by $h_{s_j,j,i}^{(-k,-\ell)}$. Put $n_{r,-k} = n_r(\{k\})$ and $n_{-k} = \sum_r n_{r,-k}$. Pool the validation contributions as

$$\begin{aligned} \widehat{\mu}_{0,\text{in}}^{(-k)} &= \frac{1}{n_{t,-k}} \sum_{\ell \neq k} \sum_{i \in \mathcal{I}_{t,\ell}} h_{0,i}^{(-k,-\ell)}, \\ \widehat{\mu}_{j,\text{in}}^{(-k)} &= \frac{1}{n_{t,-k}} \sum_{\ell \neq k} \sum_{i \in \mathcal{I}_{t,\ell}} h_{t,j,i}^{(-k,-\ell)} + \frac{1}{n_{s_j,-k}} \sum_{\ell \neq k} \sum_{i \in \mathcal{I}_{s_j,\ell}} h_{s_j,j,i}^{(-k,-\ell)}. \end{aligned} \tag{S14}$$

For each site, center its pooled validation contributions by their pooled site mean and compute their covariance matrix. Divide the target covariance by $n_{t,-k}$ and the source residual variances by $n_{s_j,-k}$ to obtain $\widehat{\Sigma}_{\text{in}}^{(-k)}$. Its off-diagonal entries retain the shared target observations. Form the discrepancies from Eq. (S14), use n_{-k} in the variance objective and $c_{n_{-k}}$ in the Wald penalty, and solve for the aggregation weights after pooling all validation contributions. For the joint TATE objective, stack both arms and use the covariance construction in Section S1.3 on these independent validation records; the alternative armwise objective in Section S7.4 instead solves separately for each arm.

Outer refits use $B = \{k\}$ in Eq. (S13). Evaluate them on fold k and apply $\widehat{\eta}^{a,(-k)}$. For balanced folds, average the outer estimates with weights $1/K_f$. Estimate final variance from the pooled sitewise outer contributions, including the two-arm covariance. Do not treat outer estimates as independent: their training samples overlap. The proof in Lemma S5 replaces the fitted contributions by common population contributions before applying the central limit theorem.

S7.2 One-Round and Two-Round Communication

In the one-round version, the target prepares messages for every $B \in \{\{k\}, \{k, \ell\} : k \neq \ell\}$ and $h \notin B$. Each message contains the initial target outcome coefficient trained outside $B \cup \{h\}$, the feature mean on that initial-training sample, and the clipped derivative moment on target calibration fold h. The source uses these quantities in Eq. (S13). It returns calibrated outcome coefficients indexed by B, plus residual means, second moments, cross-arm products, and counts on each requested validation or evaluation fold. All tasks are fixed in advance and batched into one broadcast and one return. Target predictions and their covariances are computed locally after the coefficients arrive. Joint-weight coefficients remain at the source.

In the two-round version, the first broadcast supplies all initial-training target feature means and task labels, and the first return supplies source-specific initial outcome coefficients. The second broadcast supplies the corresponding target calibration-fold derivative moments; the second return supplies the calibrated outcome coefficients and residual summaries. The dependencies are the same for inner validation and outer refits. Aggregation weights are computed locally at the target and do not require a further source refit. These round counts require a predeclared fitting protocol whose tuning and retries do not request new data-dependent target summaries.

For treatment-effect inference, each source must return the covariance of its two arm residual scores in addition to their marginal variances. If $\widehat{\xi}^{a}_{t,s_j,i}$ is the centered residual score for arm a, the required source contribution is the empirical variance of $\widehat{\eta}^{1}_{s_j}\widehat{\xi}^{1}_{t,s_j,i} - \widehat{\eta}^{0}_{s_j}\widehat{\xi}^{0}_{t,s_j,i}$. Although the uncentered arm residuals have disjoint treatment indicators, their centered covariance need not be zero. The target likewise uses the difference of its two aggregated target scores.

S7.3 Connecting Inner Validation to Outer Evaluation

Lemma S5 (Common population limits and validation moments)**.** *Use the complete calibration fits and pooled validation moments in Eqs. (S13) and (S14), with fixed $K_f \geq 4$, balanced sitewise folds, and no evaluation observations entering fitting or tuning. Use the same prespecified feature map, clipping radius, losses, and sitewise normalization for every excluded set B.*

For the target, assume the conditions of Proposition S2. For every unbiased source-assisted estimator, assume transportability and the correct-model, design, tail, and sparsity conditions of Lemma 4. Choose penalties of the stated order using each fit's own training counts and require the stated remainder conditions. For every other source-assisted estimator, assume Proposition S4 and a fixed nonzero limiting bias. Then the following statements hold simultaneously over the finitely many inner validation fits and outer refits.

(a) The calibrated coefficient limits are independent of B. The rates of Lemma S3 hold with the corresponding training sample sizes. Initial and calibrated limits are equated only for a correctly specified model.

(b) With site-centered population contributions $\varphi_t, \zeta_{t,s_j}, \xi_{t,s_j}$ as defined above, inner validation gives

$$\widehat{\mu}_{0,\mathrm{in}}^{(-k)} - \mu = \frac{1}{n_{t,-k}} \sum_{i \in \mathcal{T}_t(\{k\})} \varphi_{t,i} + o_p(n_{-k}^{-1/2}),$$
$$\widehat{\mu}_{j,\mathrm{in}}^{(-k)} - \mu = \frac{1}{n_{t,-k}} \sum_{i \in \mathcal{T}_t(\{k\})} \zeta_{t,s_j,i} + \frac{1}{n_{s_j,-k}} \sum_{i \in \mathcal{T}_{s_j}(\{k\})} \xi_{t,s_j,i} + o_p(n_{-k}^{-1/2})$$

for every unbiased source-assisted estimator. Each biased estimator instead converges to its own fixed limit.

(c) Inner and outer empirical sitewise score covariances converge to the same population covariances. On the unbiased coordinates, $n_{-k}\widehat{\Sigma}_{\mathrm{in}}^{(-k)}$ converges to the covariance used

by the full-sample oracle objective. The same statement holds jointly for the two arms. For biased coordinates, only convergence of the score-based covariance is asserted.

(d) Each outer-fold estimator has the analogous expansion on its evaluation fold. Averaging these expansions over k gives the full sitewise expansion, with each observation contributing once.

Proof of Lemma S5. Fold labels are fixed independently of the observations. Thus every sitewise empirical mean in an initial or calibration loss estimates the same site expectation, regardless of B. Its normalization is by that site's own sample size. In a correct-model branch, the correct update has the same true coefficient for every bounded clipped opposite initial fit. The other calibrated update converges to the population minimizer obtained using that true initial coefficient. For a biased candidate, both initial population minimizers are stable by Proposition S4, and their calibrated population minimizers are consequently common across B. This identifies common limiting population targets without equating misspecified initial and calibrated coefficients. Unequal sitewise fold fractions can leave the $O(N^{-1})$ conditional-gradient term in Eq. (S7); the revised rate bound absorbs it for each excluded set B.

Every initial or calibration block contains a fixed positive fraction of its site's observations. There are finitely many triples (k, ℓ, h), sources, and arms. Apply the gradient concentration, Gram concentration, and localization arguments of Lemma S3 to each block with its own initial-training sample. Allocate the failure probability over these finitely many applications and use a union bound. The penalty order and sparsity requirements are unchanged because all training sizes are proportional to the site sizes. This proves (a); independence of overlapping training sets is not needed.

For (b), condition separately on the training data for each fit excluding k, ℓ. Those data determine the calibrated coefficients, while fold ℓ remains independent of them. Write its

fitted contribution as its population-limit contribution plus their difference. The centered empirical average of this difference is $o_p(n_{-k}^{-1/2})$ by the conditional variance bound and the L_2 consistency in Eq. (S5). For an unbiased candidate, calibration bounds the uncentered population difference by $O_p(s_{\text{eff}}\lambda^2) = o_p(n_{-k}^{-1/2})$. The target uses Proposition S2. Sum the resulting bounds over the fixed number of validation folds using Eq. (S14). For biased candidates, Proposition S4 supplies consistency for their common limits; no quadratic bias bound is required.

For (c), condition on each validation fit's training data to bound the empirical squared difference between fitted and limiting contributions. Conditional Markov's inequality and L_2 consistency make this difference $o_p(1)$. Cauchy–Schwarz gives the same result for cross-products, including shared-target and cross-arm products. Laws of large numbers then apply to the common population contributions pooled within each site. Balanced folds preserve the site fractions asymptotically: $n_{r,-k}/n_{-k} = n_r/N + O(N^{-1})$. Centering and dividing by the site counts therefore gives the asserted covariance limit. Biased estimators need not have these contributions as influence functions; their convergent covariance components are used only as discrepancy scales and as a finite extension of the quadratic objective.

For (d), repeat the preceding conditional argument for $B = \{k\}$ on fold k. Each remainder is $o_p(N^{-1/2})$. For exactly equal folds, averaging assigns each observation coefficient $1/n_r$ at site r. With balanced integer fold sizes, replacing the fold weights by $1/K_f$ changes the expansion by $O_p(N^{-1})$; the same rounding bound applies when pooling inner validation folds. Using the actual site counts gives the sitewise expansion directly. The central limit theorem applies to these sitewise sums, not to independent outer-fold estimates. At no step do we condition on all overlapping fitted models and claim that all their evaluation samples remain independent. □

This lemma verifies the validation inputs needed by the aggregation theorem. It does not derive correct-model assumptions from marginal agreement: each unbiased source-assisted estimator is required to satisfy the source inference conditions. A nontransportable source whose marginal bias happens to cancel, but whose fitted model has a nonzero first derivative, is not covered by this sufficient regime.

S7.4 Armwise Covariance and Objective

For armwise aggregation, index the target-only estimator by $j = 0$ and sources by $j = 1, \ldots, K$, with $\widehat{\mu}_0 = \widehat{\mu}_t^a$, $\widehat{\mu}_j = \widehat{\mu}_{s_j}^a$, and $\eta_j = \eta_{s_j}^a$. Let μ_j^* be each candidate's probability limit, with $\mu_0^* = \mu$. Let $Z_t = (\varphi_t, \zeta_{t,s_1}, \ldots, \zeta_{t,s_K})^T$ contain the site-centered target scores. With $B = \mathrm{Cov}_t(Z_t)$, $v_j = \mathrm{Var}_{s_j}(\xi_{t,s_j})$, and $b(\eta) = (1 - \mathbf{1}^T\eta, \eta^T)^T$, the limiting quadratic variance is

$$L^*(\eta) = \frac{b(\eta)^T B b(\eta)}{\rho_t^{\mathrm{all}}} + \sum_j \frac{\eta_j^2 v_j}{\rho_{s_j}^{\mathrm{all}}}.$$

On the unbiased candidate subspace this is the asymptotic variance. For a biased candidate, its limiting fitted-score covariance may merely define a finite quadratic extension used in screening; it need not be the covariance of a valid linear expansion for that candidate.

Lemma S6 (Shared-target covariance)**.** *If all estimated score covariance components converge to finite limits, and on the unbiased subspace these are the covariances in the candidate linear expansions, the empirical version $\widehat{L}$ converges uniformly to L^* on every fixed compact weight set. Its gradient is uniformly $O_p(1)$ there.*

Proof of Lemma S6. The target sample is the only sample appearing in more than one candidate estimator, so it supplies all off-diagonal covariances. Independent source samples supply $\sum_j \eta_j^2 v_j / n_{s_j}$. With fixed K, consistency of finitely many variance and covariance components implies coefficientwise convergence of the quadratic polynomial and hence uniform convergence of it and its gradient on a compact set. The estimated quadratic

is positive semidefinite when assembled from one target empirical covariance matrix and nonnegative source variances. □

Use the main-text discrepancy and standard-error conventions, suppressing the arm index: $\widehat{d}_j = \widehat{\mu}_j - \widehat{\mu}_0$ and $\widehat{s}_j = \widehat{s}^a_{s_j}$. The variance $\widehat{s}^2_j$ includes shared-target covariance. Its scaled version $N\widehat{s}^2_j$ consistently estimates the limiting discrepancy variance for unbiased candidates; for biased candidates only a positive finite limit is required. On an outer training sample, use its size n in place of N and set $c_n = \sqrt{\log(n \vee 3)}$. Define $\widehat{t}_j = |\widehat{d}_j|/\widehat{s}_j$ and $a_{j,N} = (\widehat{t}_j/c_n - 1)_+$. The adaptive criterion is

$$\widehat{\eta} \in \arg\min_{\eta \in \mathcal{B}_\eta} \left\{ \widehat{L}(\eta) + \sum_j a_{j,N}|\eta_j| \right\}.$$

The bounded ball $\mathcal{B}_\eta = \{\eta : \|\eta\|_2 \le C_\eta\}$ contains zero and is closed under setting any coordinate to zero. These two properties will be used in the selection argument.

S7.5 Unbiased Candidates and Separation

The assumption below applies to this bounded, armwise criterion. It is distinct from the conditions for the primary joint procedure in Assumption 5. Let $\mathcal{J}^* = \{j : \mu^*_j = \mu\}$ and $b_{j,N} = |\mu^*_j - \mu|$. This is the unbiased candidate set, which need not coincide with a set of conditionally transportable populations. Its definition permits $\mathcal{J}^* = \varnothing$. Define η^* as the minimizer of L^* over the oracle subspace $\{\eta \in \mathcal{B}_\eta : \eta_j = 0 \text{ for } j \notin \mathcal{J}^*\}$.

Assumption S5 (Conditions for bounded armwise aggregation)**.** (a) The target-only estimator is asymptotically linear. Each biased candidate satisfies $\widehat{\mu}_j - \mu^*_j = o_p(1)$ for a fixed limit $\mu^*_j \neq \mu$, including on all training samples used to construct weights. Each unbiased candidate has the source-assisted linear expansion. All estimated covariance components converge, and $\sqrt{N}\widehat{s}_j \to_p s_j \in (0, \infty)$.

(b) If $\mathcal{J}^* \neq \varnothing$, the oracle quadratic is strictly convex on its subspace and its unique minimizer is interior to the ball. If $\mathcal{J}^* = \varnothing$, the oracle subspace is $\{0\}$ and $\eta^* = 0$ without a curvature requirement.

(c) If there are biased candidates, put $b_{0,N} = \min_{j \notin \mathcal{J}^*} b_{j,N}$. Require $\inf_N b_{0,N} > 0$, $c_n \to \infty$, and $c_n/\sqrt{n} \to 0$ on each outer training sample. If all candidates are unbiased, only $c_n \to \infty$ is needed.

(d) The number of sites and folds is fixed, all site fractions are positive, and the stated moment bounds and positive limiting variances hold for the final oracle aggregate.

The requirement on biased candidates is only consistency and a positive finite limiting discrepancy scale. Proposition S4 gives a sufficient route from stable sparse initial fits to these properties for the specified GLM class. Orthogonality or root-N inference for those candidates is unnecessary. For a fixed nonzero bias, consistency implies $|\widehat{d}_j|/b_{j,N} \to_p 1$, whereas the unbiased-candidate expansions imply $\sqrt{N}\widehat{d}_j = O_p(1)$. Consequently,

$$\widehat{t}_j = O_p(1) \quad (j \in \mathcal{J}^*), \qquad \widehat{t}_j = \Theta_p(\sqrt{N} b_{j,N}) \quad (j \notin \mathcal{J}^*).$$

The proposed growing cutoff satisfies the threshold window for these fixed biases. For a drifting bias, a separate sufficient extension requires candidate errors $o_p(b_{0,N})$ and $c_n/(\sqrt{n} b_{0,N}) \to 0$, in addition to the other uniform conditions; consistency alone is insufficient in that regime. The retained fixed cutoff 2 lies outside the theorem's growing-cutoff regime.

S7.6 Oracle Aggregation Result

Theorem S1 (Adaptive aggregation limit)**.** *Under Assumption S5 and the target and source expansion conditions above,* $\|\widehat{\eta} - \eta^*\|_2 = o_p(1)$ *and*

$$P\{\widehat{\eta}_j = 0 \text{ for all } j \notin \mathcal{J}^*\} \longrightarrow 1.$$

The outer-cross-fitted aggregate has the same first-order distribution as the oracle:

$$\sqrt{N}(\widehat{\mu}_{\text{agg}} - \mu) \xrightarrow{d} \mathcal{N}\{0, L^*(\eta^*)\}, \qquad \widehat{L}(\widehat{\eta}) \to_p L^*(\eta^*).$$

If no candidate is unbiased, selection returns the target-only estimator with probability tending to one.

S7.7 Selection and Oracle Equivalence

Proof of Theorem S1. For an unbiased candidate, $\widehat{t}_j/c_n \to_p 0$, so its penalty is exactly zero with probability tending to one. For a biased candidate, $a_{j,N} \to_p \infty$ by separation. Fixed K makes these statements simultaneous over all coordinates.

Write $G_N = \sup_{\eta \in \mathcal{B}_\eta} \|\nabla \widehat{L}(\eta)\|_\infty = O_p(1)$. For any η with $\eta_j \neq 0$, let $\eta^{(j=0)}$ be obtained by setting that coordinate to zero. It remains in the ball, and the mean-value inequality gives

$$\widehat{L}(\eta) - \widehat{L}(\eta^{(j=0)}) \geq -G_N|\eta_j|.$$

All other penalty terms are unchanged. Thus the full objective at η exceeds its value at $\eta^{(j=0)}$ whenever $a_{j,N} > G_N$. This event holds for every biased coordinate with probability tending to one, so every minimizer excludes them. This argument also covers $\mathcal{J}^* = \varnothing$ directly.

On the oracle subspace, all remaining penalties vanish with probability tending to one. Uniform convergence of $\widehat{L}$ and uniqueness of its population minimizer imply $\widehat{\eta} \to_p \eta^*$ by the

compact argmin argument. The same reasoning applies to each outer-training fit; fixed fold count makes the conclusion simultaneous.

For outer fold k, the difference between the fitted-weight and oracle-weight estimates is

$$\sum_{j \in \mathcal{J}^*} (\widehat{\eta}_j^{(-k)} - \eta_j^*)(\widehat{\mu}_{j,k} - \widehat{\mu}_{0,k}) = o_p(N^{-1/2}),$$

because each coefficient difference is $o_p(1)$ and each unbiased candidate difference is $O_p(N^{-1/2})$. Averaging over finitely many folds gives oracle equivalence. Replacing the candidate estimates by their linear expansions yields a sum of independent, site-centered oracle scores. The sitewise central limit theorem proves the normal limit. Finally, uniform covariance consistency and $\widehat{\eta} \to_p \eta^*$ imply $\widehat{L}(\widehat{\eta}) \to_p L^*(\eta^*)$, and Slutsky's theorem gives studentized inference. □

The oracle arm-mean variance is bounded above by $L^*(0)$ because the target-only weight vector is feasible. A strict gain requires a nonzero source direction that lowers the quadratic. No assumption that at least one source exists or receives positive weight is required for validity. Separate minimization in the two arms does not minimize the variance of their difference; its covariance terms must be retained as described in the main text.

S7.8 Squared Discrepancies and Local Bias

The raw-discrepancy version of the FACE criterion (Han et al., 2025) replaces $a_{j,N}$ by $\lambda_{\mathrm{sq},N}\widehat{d}_j^2$. On unbiased coordinates its penalty is $O_p(\lambda_{\mathrm{sq},N}/N)$; on separated biased coordinates it has order $\lambda_{\mathrm{sq},N}b_{j,N}^2$. Sufficient conditions for the same compact-set selection and consistency argument are

$$\lambda_{\mathrm{sq},N}/N \to 0, \qquad \lambda_{\mathrm{sq},N}b_{0,N}^2 \to \infty.$$

The more restrictive choice $\lambda_{\mathrm{sq},N} = o(\sqrt{N})$ also makes the unbiased-coordinate penalty $o_p(N^{-1/2})$. That stronger restriction is not needed merely for consistency of weights in

the cross-fitted oracle-equivalence proof above. The Wald form is used here because it places discrepancies from differently variable sources on a comparable scale, not because an intrinsically different detection boundary follows without specifying the tuning regime.

If $b_{j,N} = O(N^{-1/2})$, the biased Wald statistic need not diverge. Exact selection is then not guaranteed. For a fixed weight vector the first-order mean shift is $\sqrt{N}\sum_j \eta_j(\mu_j^* - \mu)$. Biases of smaller order are negligible, but biases on the exact root-N scale may shift the limit. If data-adaptive weights fail to converge to a deterministic limit, even this fixed-weight normal approximation does not automatically justify inference. Addressing such local alternatives requires an additional analysis.

References

Bang, H. and Robins, J. M. (2005), 'Doubly robust estimation in missing data and causal inference models', *Biometrics* **61**(4), 962–973.

Connors, A. F., Speroff, T., Dawson, N. V., Thomas, C., Harrell, F. E., Wagner, D., Desbiens, N., Goldman, L., Wu, A. W., Califf, R. M. et al. (1996), 'The effectiveness of right heart catheterization in the initial care of critically ill patients', *JAMA* **276**(11), 889–897.

Han, L., Hou, J., Cho, K., Duan, R. and Cai, T. (2025), 'Federated adaptive causal estimation (FACE) of target treatment effects', *Journal of the American Statistical Association* **120**(551), 1503–1516.

Hirano, K. and Imbens, G. W. (2001), 'Estimation of causal effects using propensity score weighting: An application to data on right heart catheterization', *Health Services and Outcomes Research Methodology* **2**(3), 259–278.

Hou, J., Mukherjee, R. and Cai, T. (2025), 'Efficient and robust semi-supervised estimation of average treatment effect with partially annotated treatment and response', *Journal of Machine Learning Research* **26**(40), 1–77.

Negahban, S. N., Ravikumar, P., Wainwright, M. J. and Yu, B. (2012), 'A unified framework for high-dimensional analysis of m-estimators with decomposable regularizers', *Statistical Science* **27**(4), 538–557.

Smucler, E., Rotnitzky, A. and Robins, J. M. (2019), 'A unifying approach for doubly-robust ℓ_1 regularized estimation of causal contrasts', *arXiv preprint arXiv:1904.03737* .

Tan, Z. (2020), 'Model-assisted inference for treatment effects using regularized calibrated estimation with high-dimensional data', *The Annals of Statistics* **48**(2), 811–837.

Xiong, R., Koenecke, A., Powell, M., Shen, Z., Vogelstein, J. T. and Athey, S. (2023), 'Federated causal inference in heterogeneous observational data', *Statistics in Medicine* **42**(24), 4418–4439.